

Effective Actors

Jamie Allen



Who Am I?

- Consultant at Typesafe
- Actor & Scala developer since 2009

jamie.allen@typesafe.com
@jamie_allen

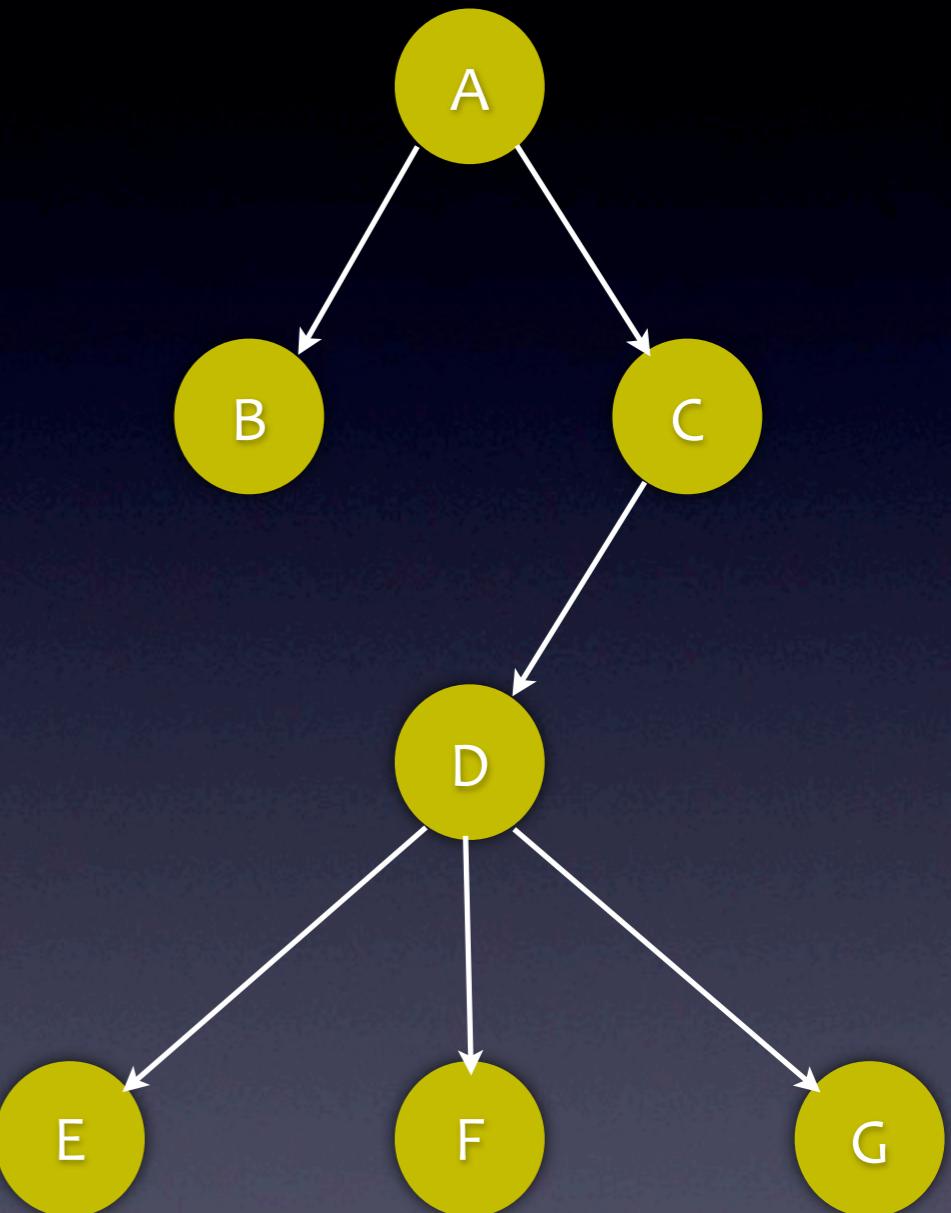


Effective Actors

- Best practices based on several years of actor development
- Helpful hints for reasoning about actors at runtime

Actors

- Concurrent, lightweight processes that communicate through asynchronous message passing
- Isolation of state, no internal concurrency



Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case _ => println("Pinging!"); sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case _ => println("Ponging!"); sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Ping!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case _ => println("Pinging!"); sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case _ => println("Ponging!"); sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Ping!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case _ => println("Pinging!"); sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case _ => println("Ponging!"); sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Ping!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case _ => println("Pinging!"); sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case _ => println("Ponging!"); sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Ping!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```



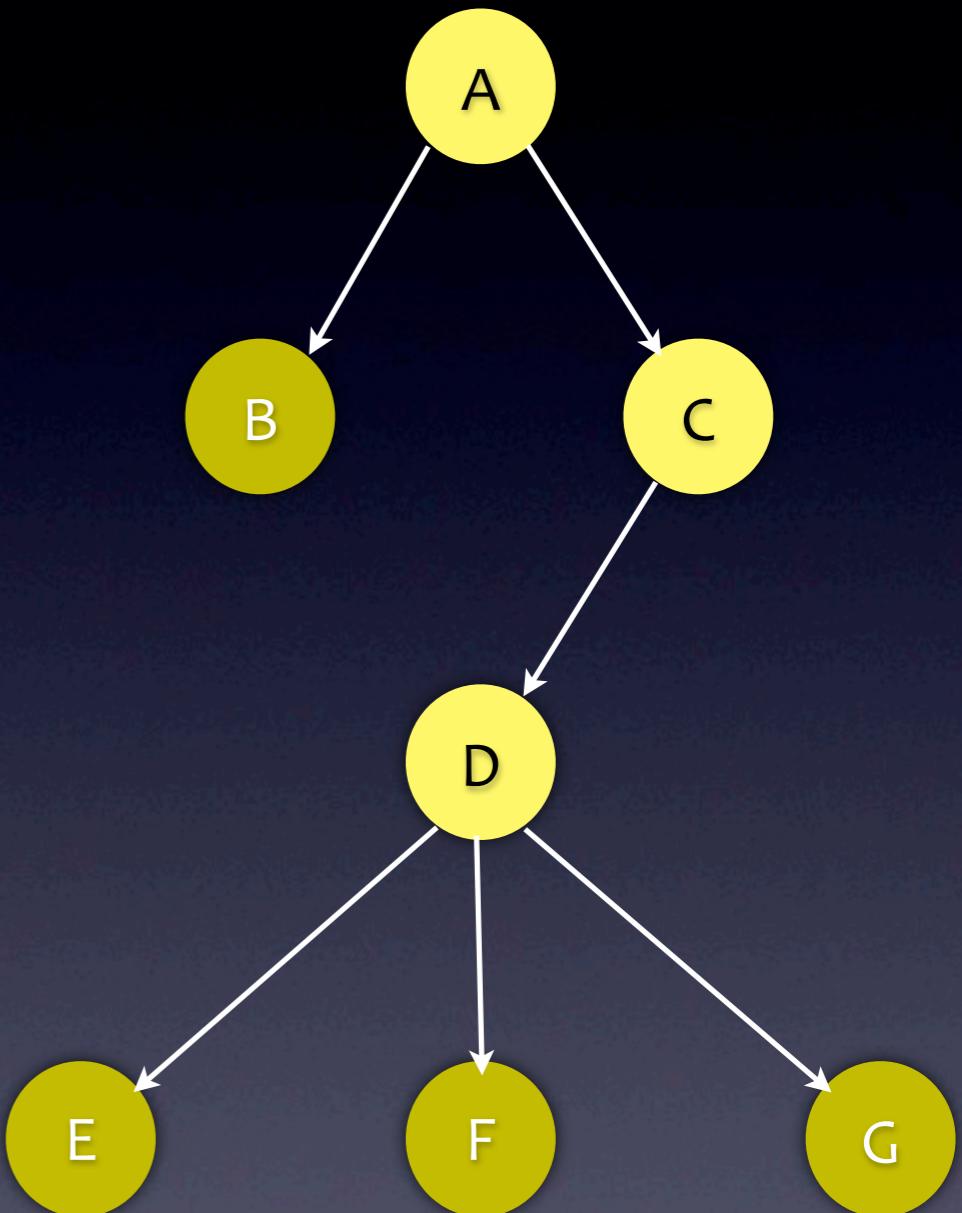
Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case _ => println("Pinging!"); sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case _ => println("Ponging!"); sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Ping!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Supervisor Hierarchies

- Specifies handling mechanisms for groupings of actors in parent/child relationship



Akka Supervisors

```
class MySupervisor extends Actor {  
    override val supervisorStrategy =  
        OneForOneStrategy() {  
            case ae: ArithmeticException => Resume  
            case np: NullPointerException => Restart  
        }  
  
    context.actorOf(Props[MyActor])  
}
```



Akka Supervisors

```
class MySupervisor extends Actor {  
    override val supervisorStrategy =  
        OneForOneStrategy() {  
            case ae: ArithmeticException => Resume  
            case np: NullPointerException => Restart  
        }  
  
    context.actorOf(Props[MyActor])  
}
```



Akka Supervisors

```
class MySupervisor extends Actor {  
    override val supervisorStrategy =  
        OneForOneStrategy() {  
            case ae: ArithmeticException => Resume  
            case np: NullPointerException => Restart  
        }  
  
    context.actorOf(Props[MyActor])  
}
```



Akka Supervisors

```
class MySupervisor extends Actor {  
    override val supervisorStrategy =  
        OneForOneStrategy() {  
            case ae: ArithmeticException => Resume  
            case np: NullPointerException => Restart  
        }  
  
    context.actorOf(Props[MyActor])  
}
```



Domain Supervision

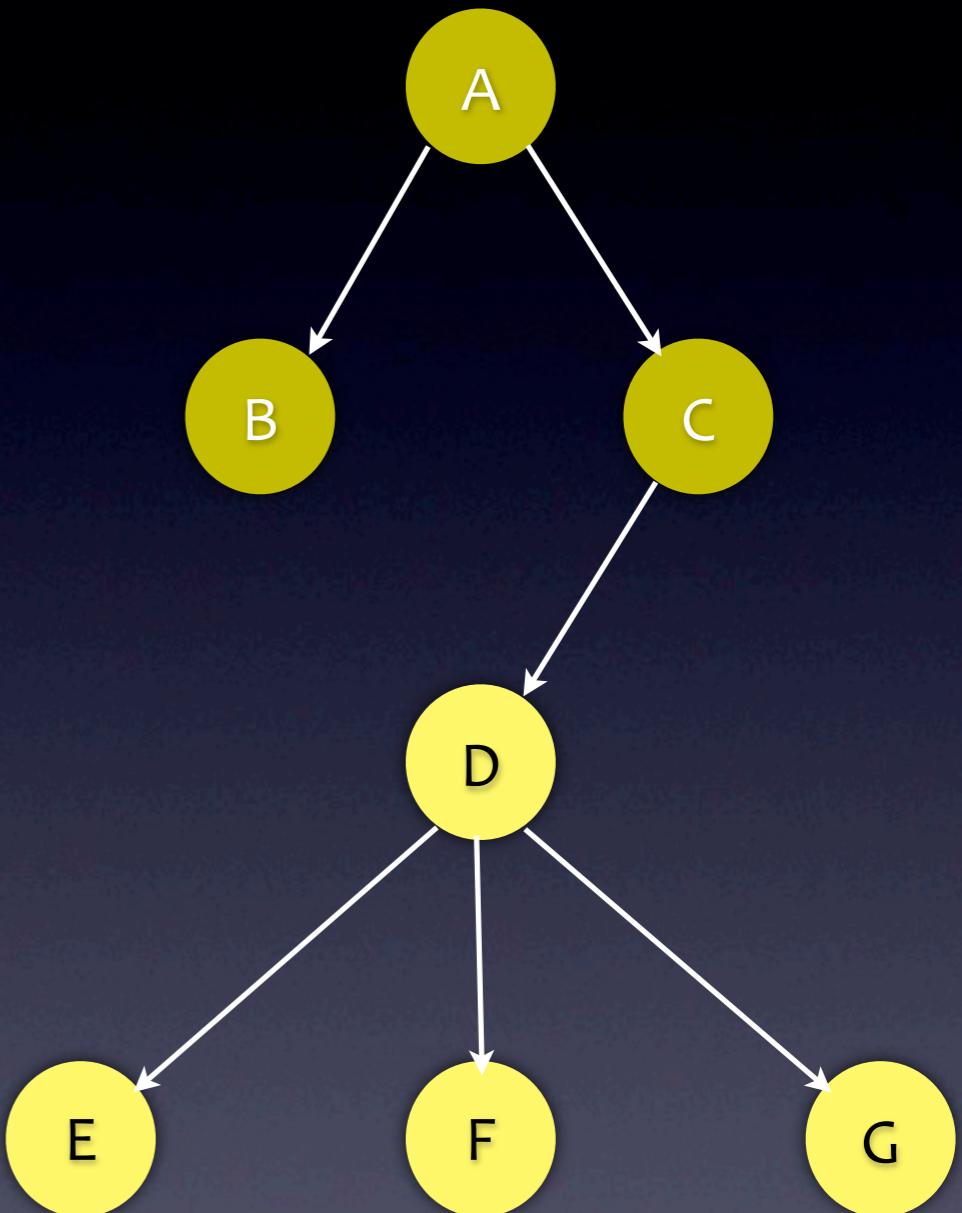
- Each supervisor manages a grouping of types in a domain
- Actors persist to represent existence of instances
- Actors constantly resolve the world as it should be against the world as it is

Worker Supervision

- Supervisors should hold all critical data
- Workers should receive data for tasks in messages
- Workers being supervised should perform dangerous tasks
- Supervisor should know how to handle failures in workers in order to retry appropriately

Parallelism

- Easily scale a task by creating multiple instances of an actor and applying work using various strategies
- Order is not guaranteed, nor should it be



Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```



Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```



Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```



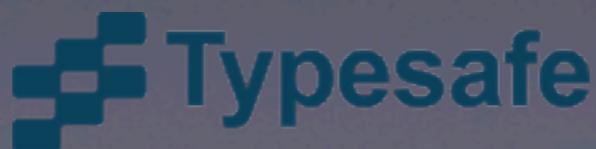
Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```

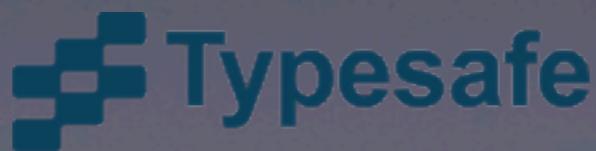


Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```



RULE: Actors Should Only Do One Thing



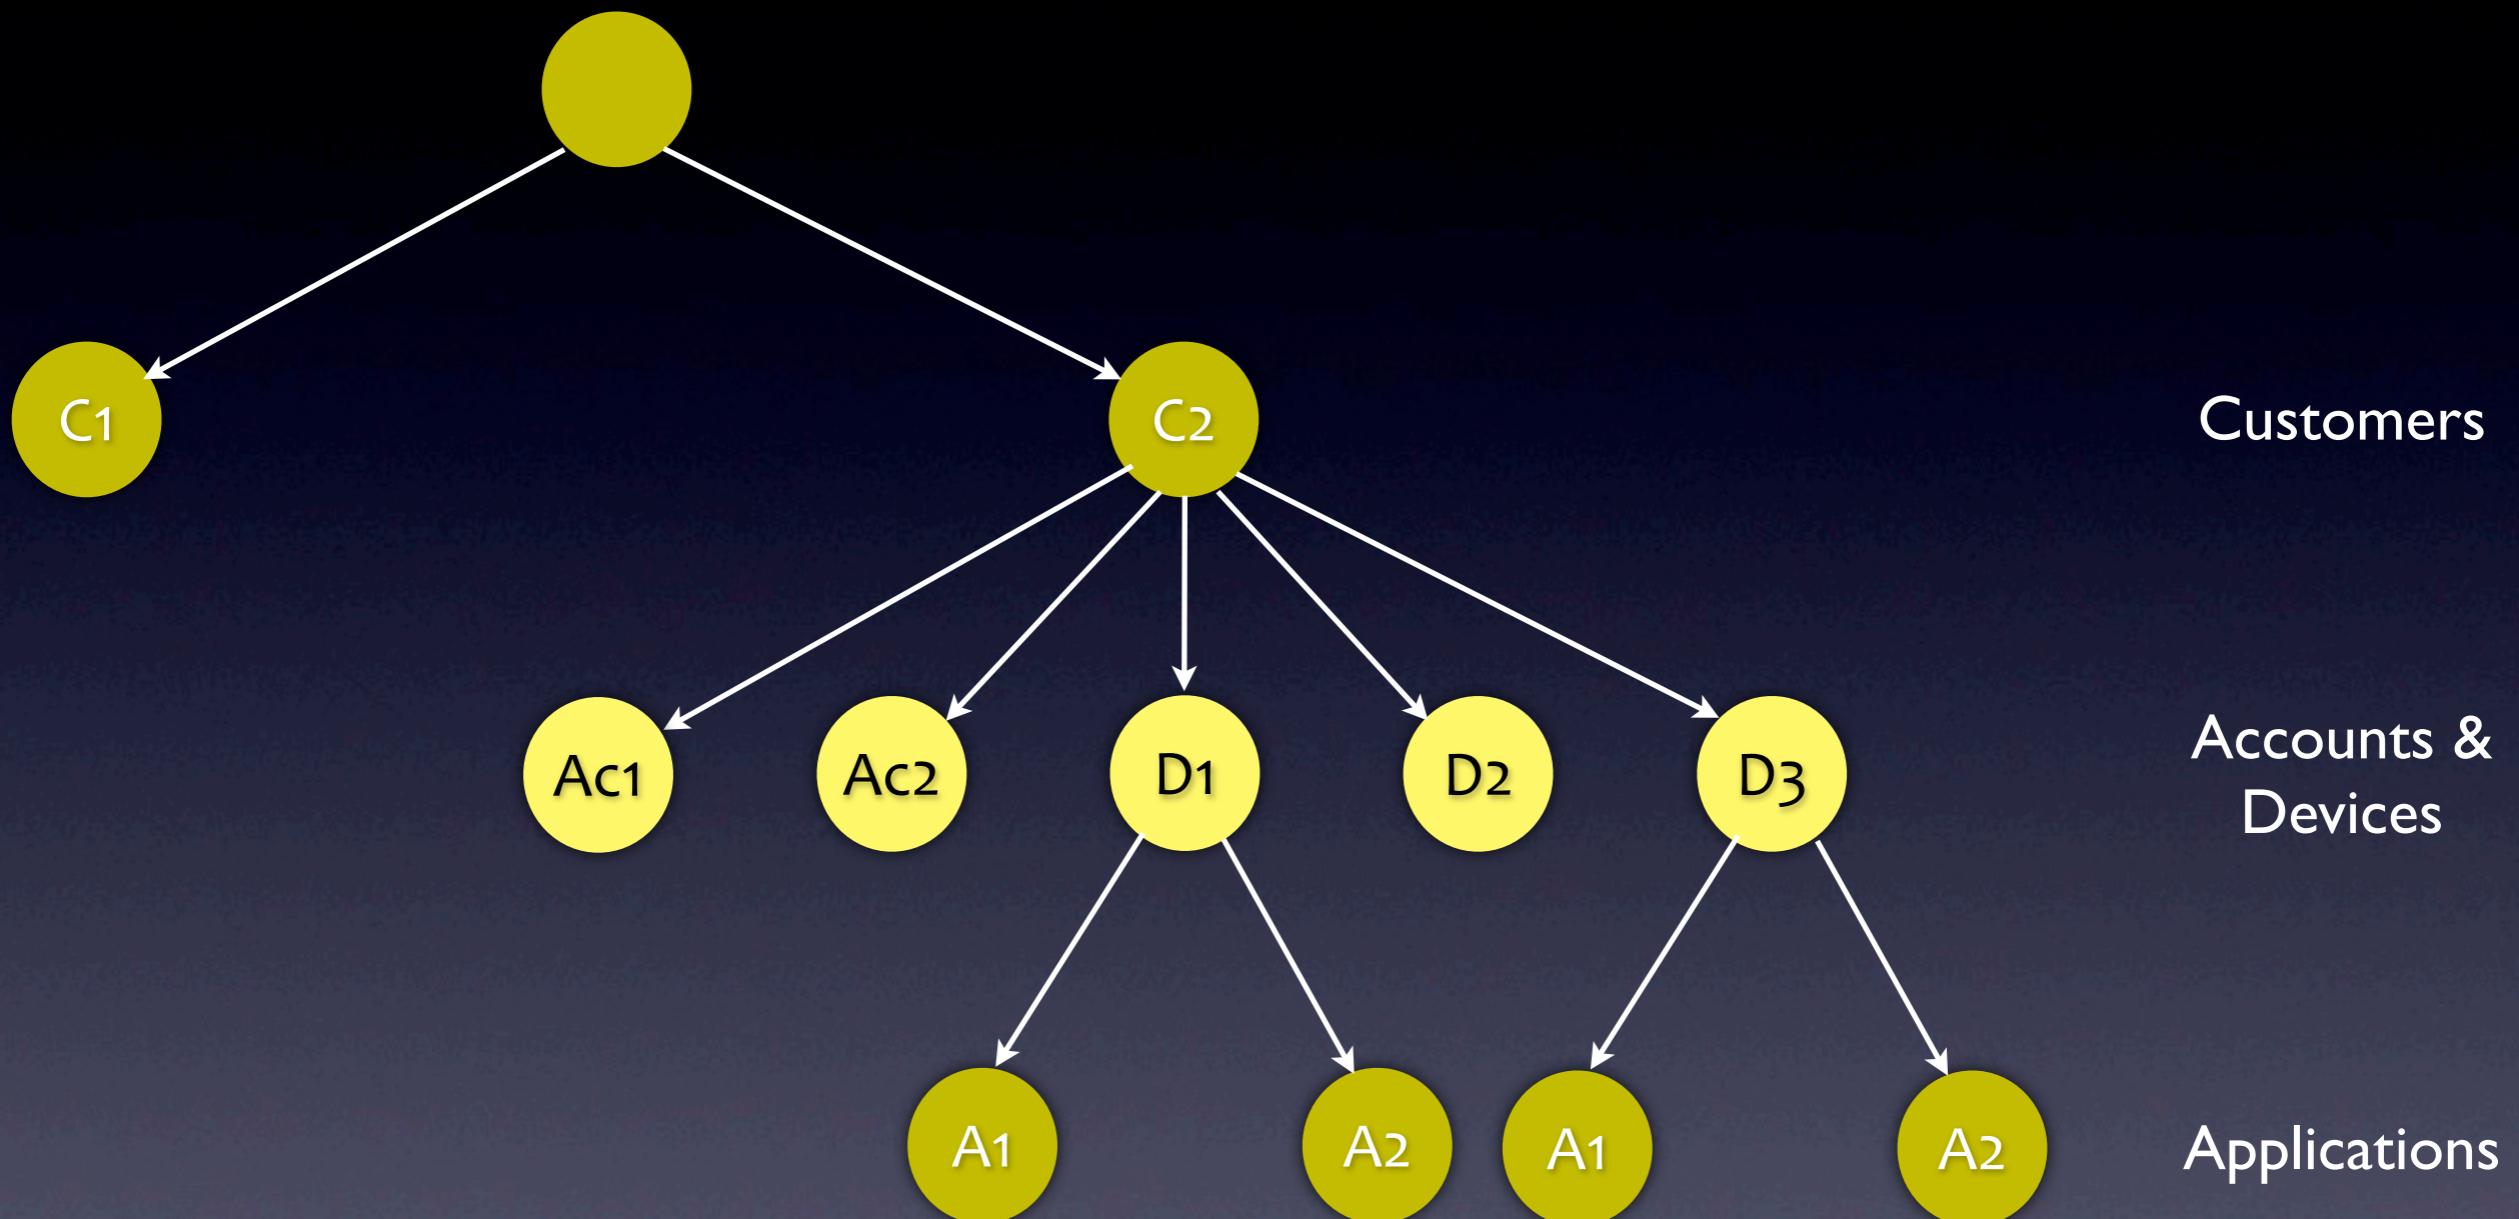
Single Responsibility Principle

- Do not conflate responsibilities in actors
- Becomes hard to define the boundaries of responsibility
- Supervision becomes more difficult as you handle more possibilities
- Debugging becomes very difficult

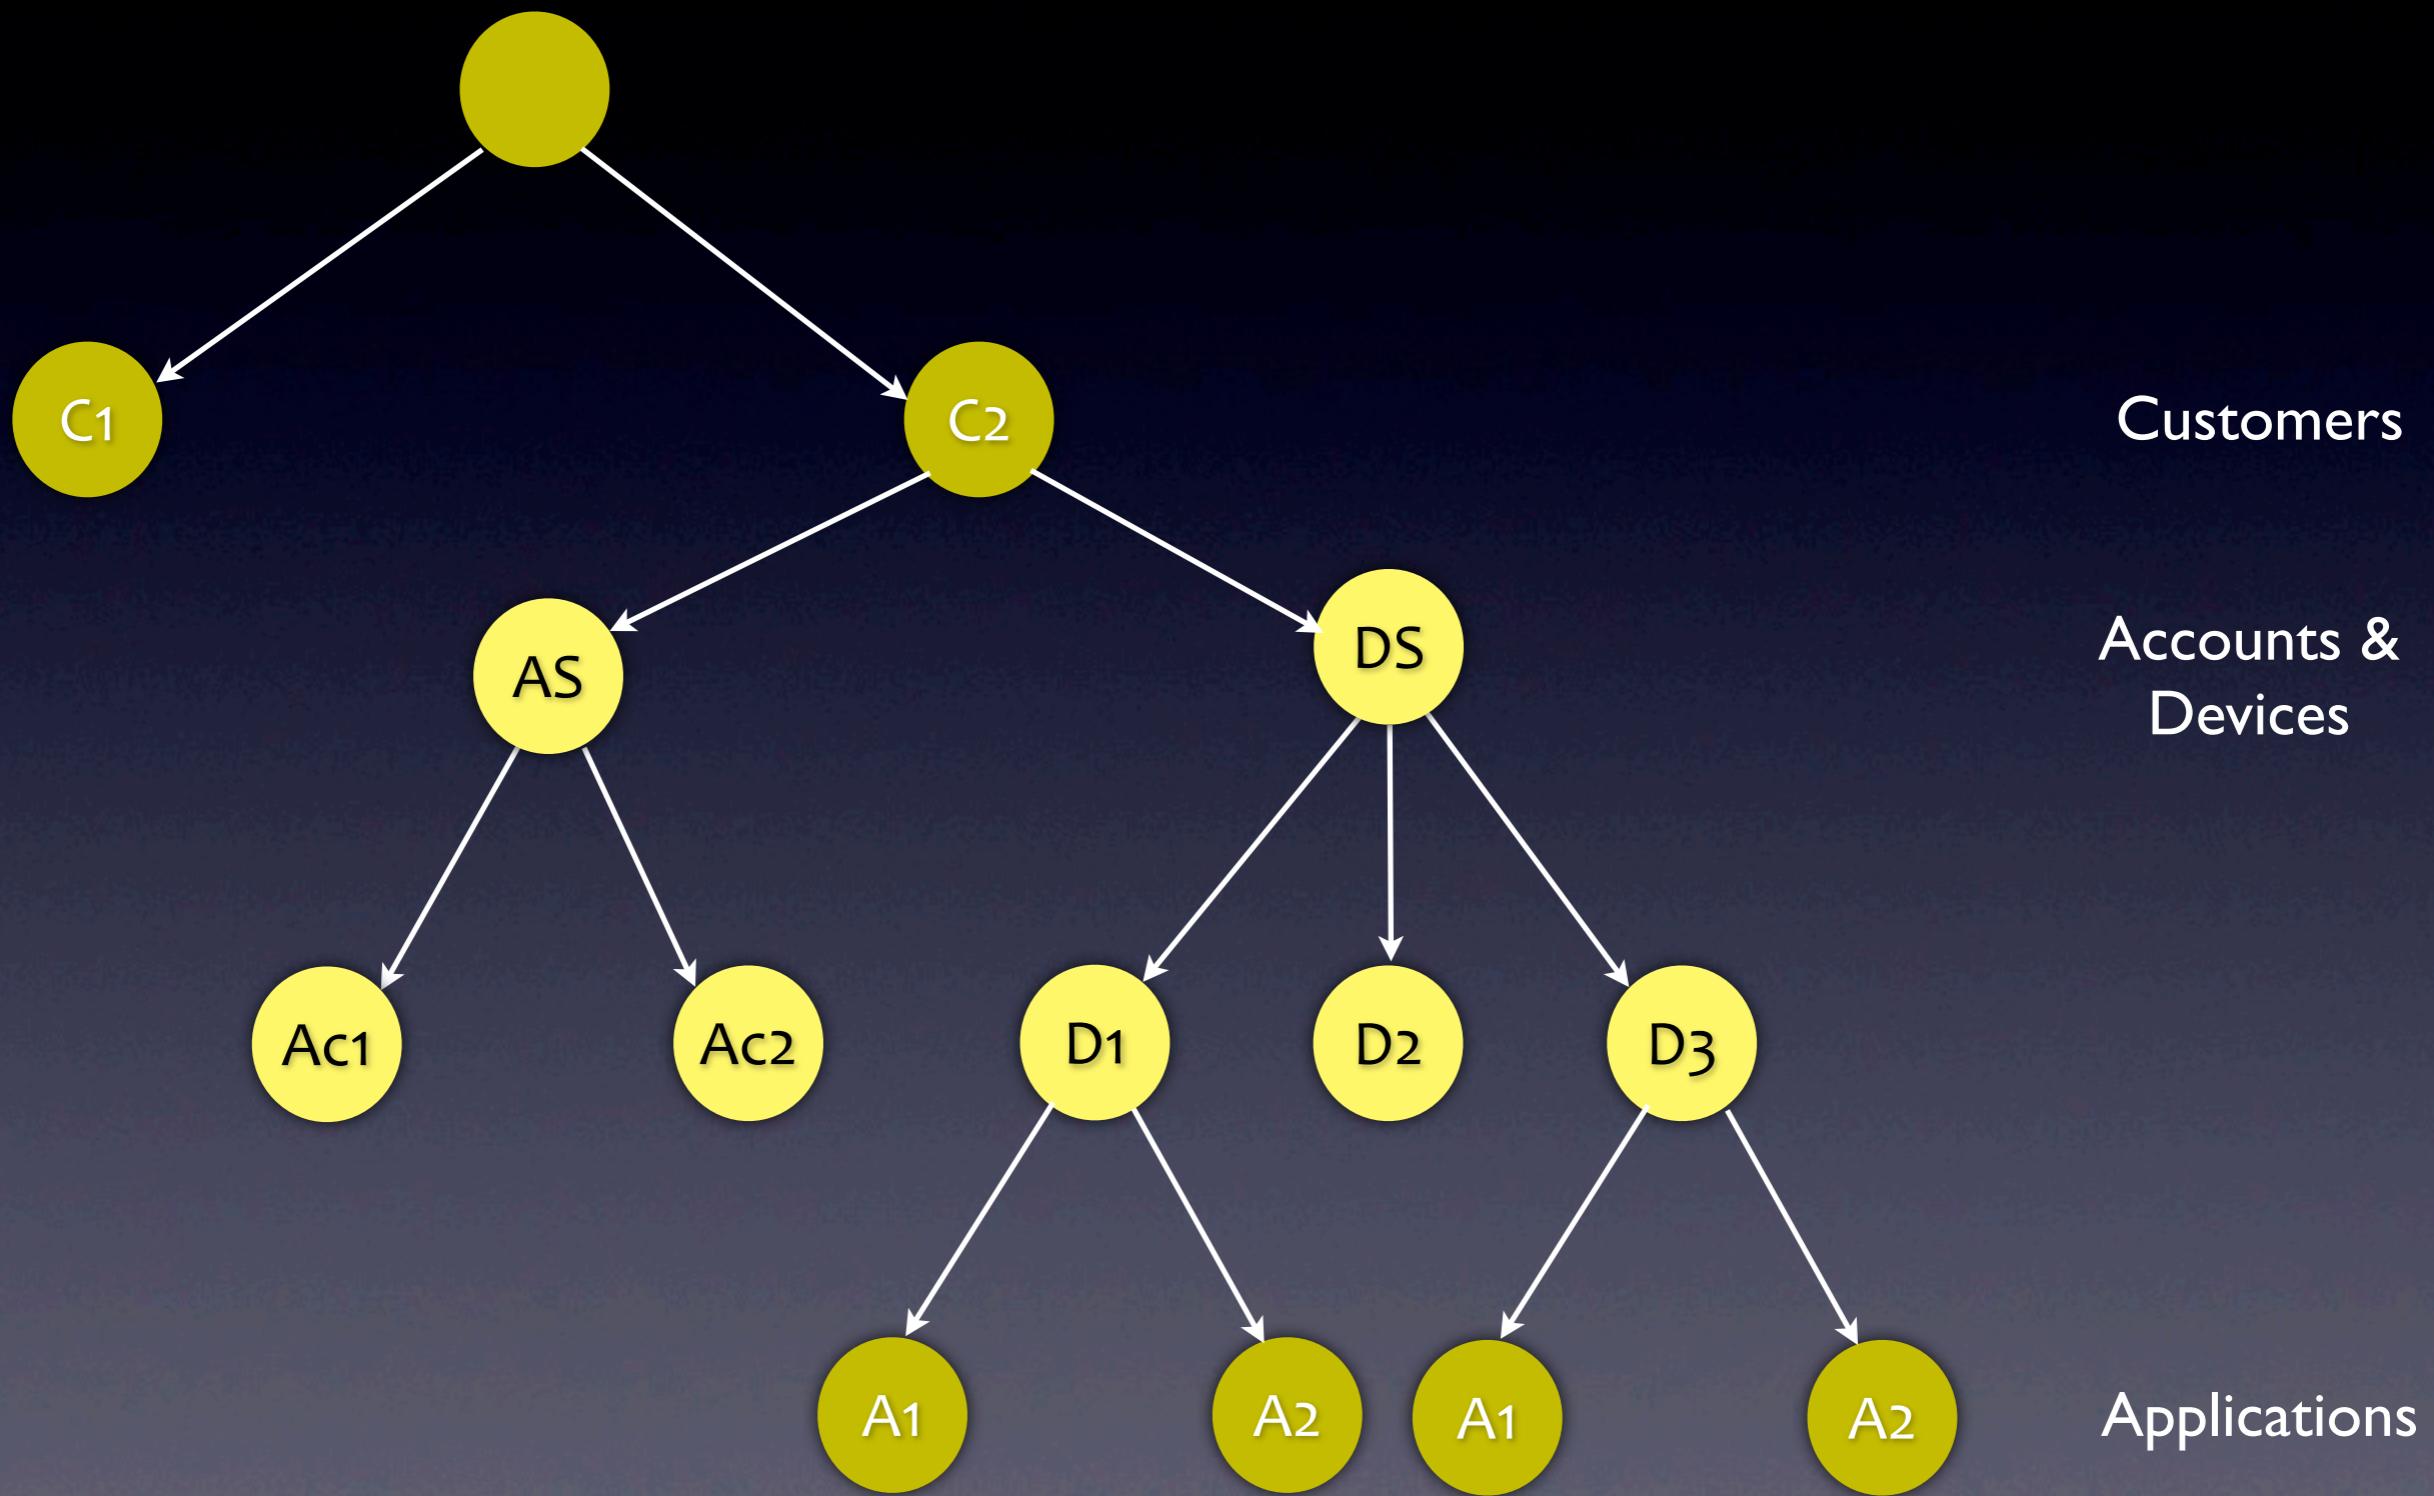
Supervision

- Every non-leaf node is technically a supervisor
- Create explicit supervisors under each node for each type of child to be managed

Conflated Supervision



Explicit Supervision



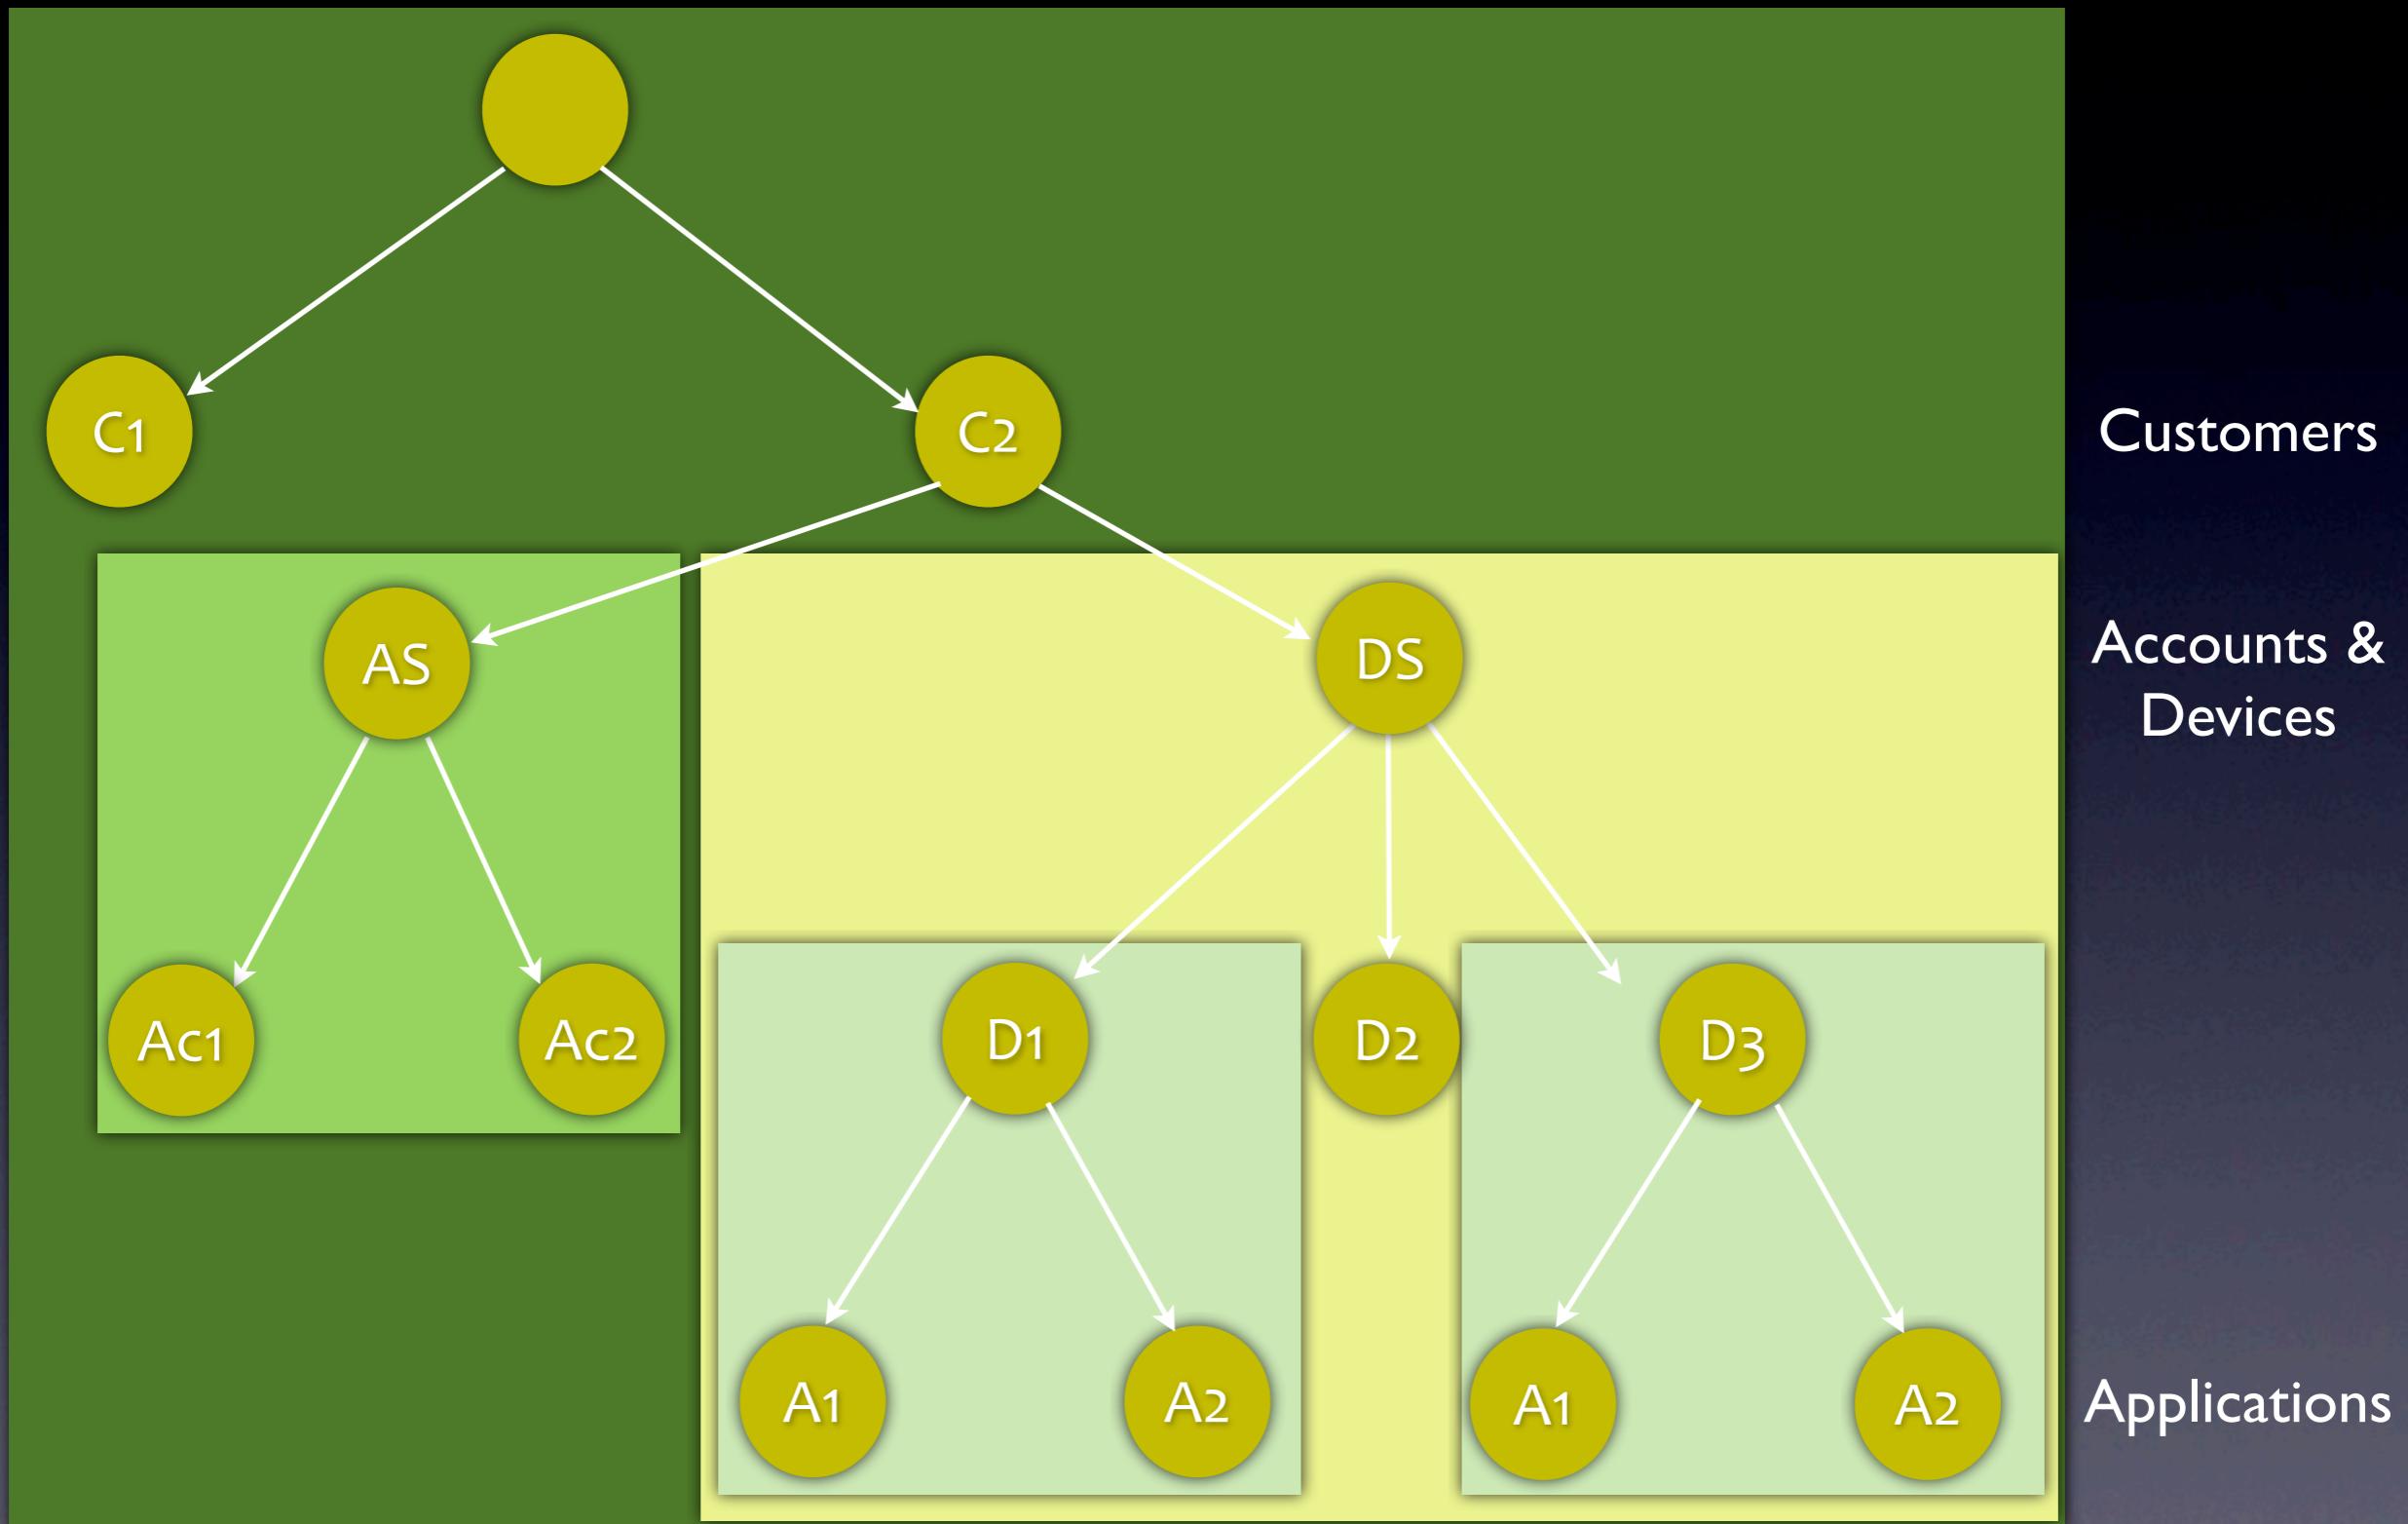
Keep the Error Kernel Simple

- Limit the number of supervisors you create at this level
- Helps with fault tolerance and explicit handling of errors through the hierarchy
- Akka uses synchronous messaging to create top-level actors

Use Failure Zones

- Multiple isolated zones with their own resources (thread pools, etc)
- Prevents starvation of actors
- Prevents issues in one branch from affecting another

Failure Zones



Takeaway

- Isolation of groups of actors limits the effects of failure
- For reasonably complex actor systems, shallow trees are a smell test
- Actors are cheap - use them

RULE: Block Only When and Where You Must



Consequences of Blocking

- Eventually results in actor starvation as thread pool dries up
- Horrible performance
- Massive waste of system resources

Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case _ =>  
        worker ? (1 to 100) map  
            (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case _ =>  
        worker ? (1 to 100) map  
            (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case _ =>  
        worker ? (1 to 100) map  
            (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case _ =>  
        worker ? (1 to 100) map  
            (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case _ =>  
        worker ? (1 to 100) map  
            (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case _ =>  
        worker ? (1 to 100) map  
            (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```



Sequential Futures

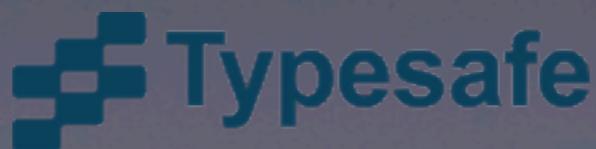
```
val r: Future[Int] = for {
  a <- (service1 ? GetResult).mapTo[Int]
  b <- (service2 ? GetResult(a)).mapTo[Int]
} yield a * b
```



Parallel Futures

First way: Define futures first

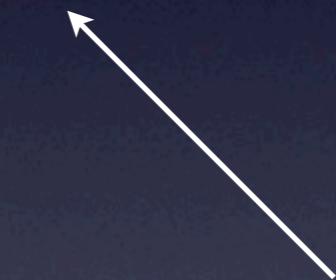
```
val f1: Future[Int] = service1 ? GetResult  
val f2: Future[Int] = service2 ? GetResult  
  
val r: Future[Int] = for {  
    a <- f1  
    b <- f2  
} yield a * b
```



Parallel Futures

Second way: Define in one line of for comprehension

```
val r: Future[Int] = for {  
    (a: Int, b: Int) <- (service1 ? GetResult) zip (service2 ? GetResult)  
} yield a * b
```



Note the “zip”

Blocking

- An example is database access
- Use a specialized actor with its own resources
- Pass messages to other actors to handle the result
- Akka provides “Managed Blocking” to limit the number of blocking operations taking place in actors at one time

Akka Dispatcher

```
class Worker extends Actor {
  def receive = {
    case s: Seq[Int] => sender ! s.reduce(_ + _)
  }
}

class Delegator extends Actor {
  implicit val timeout: Timeout = 2 seconds
  val worker = context.actorOf(Props[Worker]).withDispatcher("my-dispatcher")
  def receive = {
    case _ =>
      blocking {
        val futResult = worker ? (1 to 100)
        val result = Await.result(futResult, 2 seconds)
      }
  }
}
```



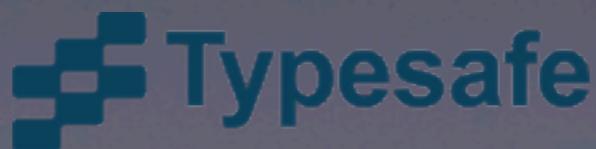
Push, Not Pull

- Start with no guarantees about delivery
- Add guarantees only where you need them
- Retry until you get the answer you expect
- Switch your actor to a "nominal" state at that point

Takeaway

- Find ways to ensure that your actors remain asynchronous and non-blocking
- Avoid making your actors wait for anything while handling a message

RULE: Do Not Optimize Prematurely



Start Simple

- Make Donald Knuth happy
- Start with a simple configuration and profile
- Do not parallelize until you know you need to and where



Initial Focus

- Deterministic
- Declarative
- Immutable
- Start with functional programming and go from there

Advice From Jonas Bonér

- Layer in complexity
- Add indeterminism
- Add mutability in hot spots
(CAS and STM)
- Add explicit locking and
threads as a last resort



Photo courtesy of Brian Clapper, NE Scala 2011

Prepare for Race Conditions

- Write actor code to be agnostic of time and order
- Actors should only care about now, not that something happened before it
- Actors can "become" or represent state machines to represent transitions

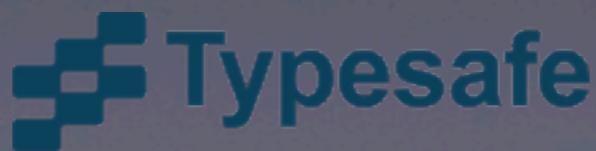
Beware the Thundering Herd

- Actor systems can be overwhelmed by "storms" of messages flying about
- Do not pass generic messages that apply to many actors
- Dampen actor messages if the exact same message is being handled repeatedly within a certain timeframe
- Tune your dispatchers and mailboxes via back-off policies and queue sizes
- Akka has added Circuit Breakers to v2.1

Takeaway

- Start by thinking of terms of an implementation that is deterministic and not actor based
- Layer in complexity as you go

RULE: Be Explicit In Your Intent



Name Your Actors

- Allows for external configuration
- Allows for lookup
- Better semantic logging

Create Specialized Messages

- Non-specific messages about general events are dangerous

`AccountsUpdated`

- Can result in "event storms" as all actors react to them
- Use specific messages forwarded to actors for handling

`AccountDeviceAdded (acctNum, deviceNum)`

Create Specialized Exceptions

- Don't use Exception to represent failure in an actor
- Specific exceptions can be handled explicitly
- State can be transferred between actor incarnations in Akka (if need be)

Takeaway

- Be specific in everything you do
- Makes everything that occurs in your actor system more clear to other developers maintaining the code
- Makes everything more clear in production

RULE: Do Not Expose Your Actors



No Direct References

- Actors die
- Doesn't prevent someone from calling into an actor with another thread
- Akka solves this with the ActorRef abstraction
- Erlang solves this with PIDs

Never Publish “this”

- Don't send it anywhere
- Don't register it anywhere
- Particularly with future callbacks
- Publish “self” instead, which is an ActorRef
- Avoid closing over "sender" in Akka, it will change with the next message

Use Immutable Messages

- Enforces which actor owns the data
- If mutable state can escape, what is the point of using an actor?

Pass Copies of Mutable Data

- Mutable data in actors is fine
- But data can escape your scope
- Copy the data and pass that, as Erlang does (COW)
- Akka has STM references

Avoid Sending Behavior

- Unless using Agents, of course
- Closures make this possible (and easy)
- Also makes it easy for state to escape

Takeaway

- Keep everything about an actor internal to that actor
- Be vary wary of data passed in closures to anyone else

RULE: Make Debugging Easy On Yourself



Externalize Business Logic

- Consider using external functions to encapsulate complex business logic
- Easier to unit test outside of actor context
- Not a rule of thumb, but something to consider as complexity increases
- Not as big of an issue with Akka's TestKit

Use Semantically Useful Logging

- Trace-level logs should have output that you can read easily
- Use line-breaks and indentation
- Both Akka and Erlang support hooking in multiple listeners to the event log stream

Unique IDs for Messages

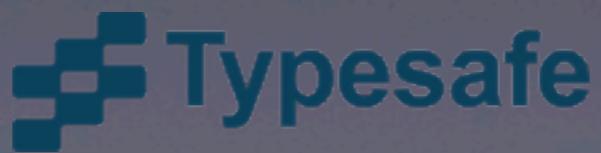
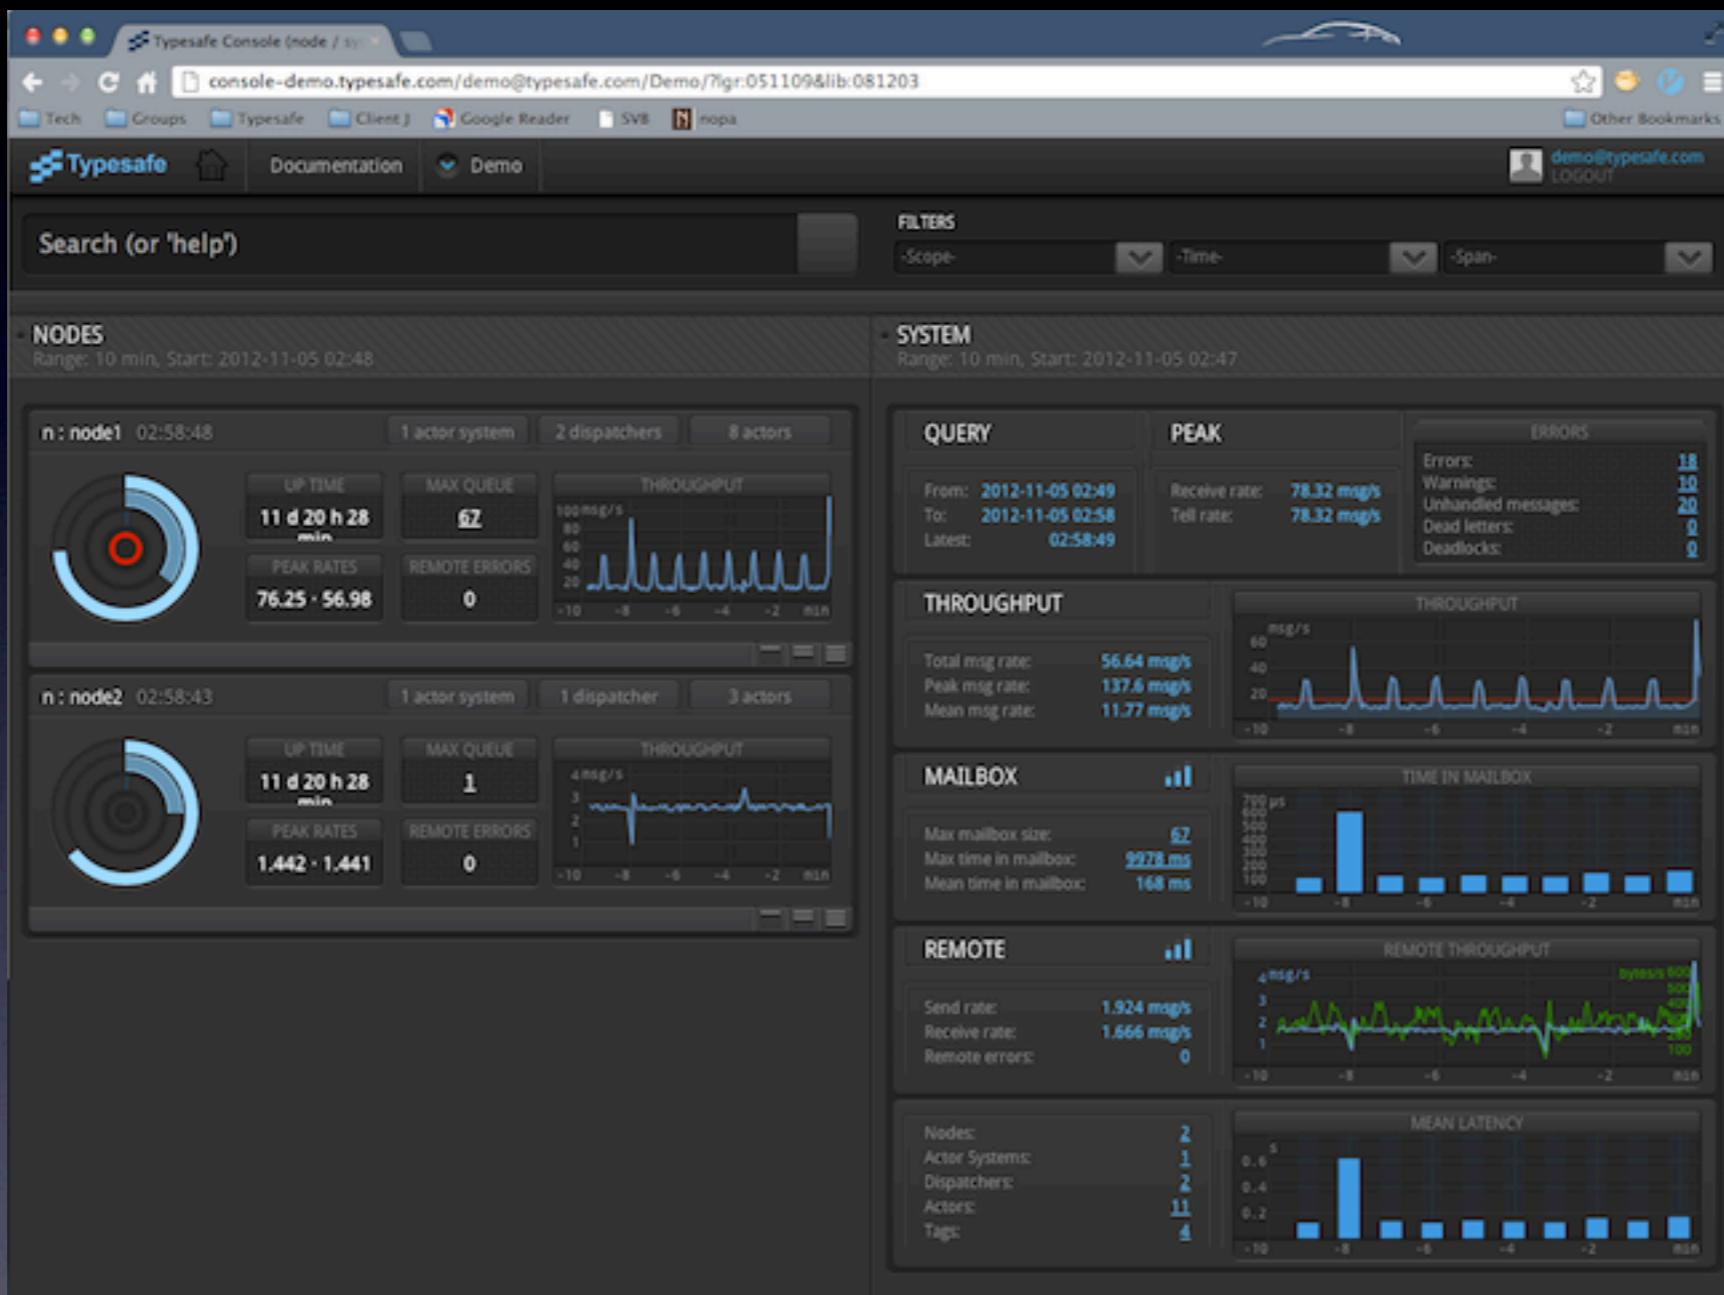
- Allows you to track message flow
- When you find a problem, get the ID of the message that led to it
- Use the ID to grep your logs and display output just for that message flow
- Akka ensures ordering on a per actor basis, also in logging

Monitor Everything

- Do it from the start
- Use tools like JMX MBeans to visualize actor realization
- The Atmos/Typesafe Console is a great tool to visualize actor systems, doesn't require you to do anything up front
- Visual representations of actor systems at runtime are invaluable



Typesafe Console



Takeaway

- Build your actor system to be maintainable from the outset
- Utilize all of the tools at your disposal

Thank You!

- Some content provided by members of the Typesafe team, including:
 - Jonas Bonér
 - Viktor Klang
 - Roland Kuhn
 - Havoc Pennington

