

# Effective Actors

Jamie Allen



# Who Am I?

- Consultant at Typesafe
- Actor & Scala developer since 2009

jamie.allen@typesafe.com

@jamie\_allen

Typesafe 社でコンサルタントしています

# Effective Actors

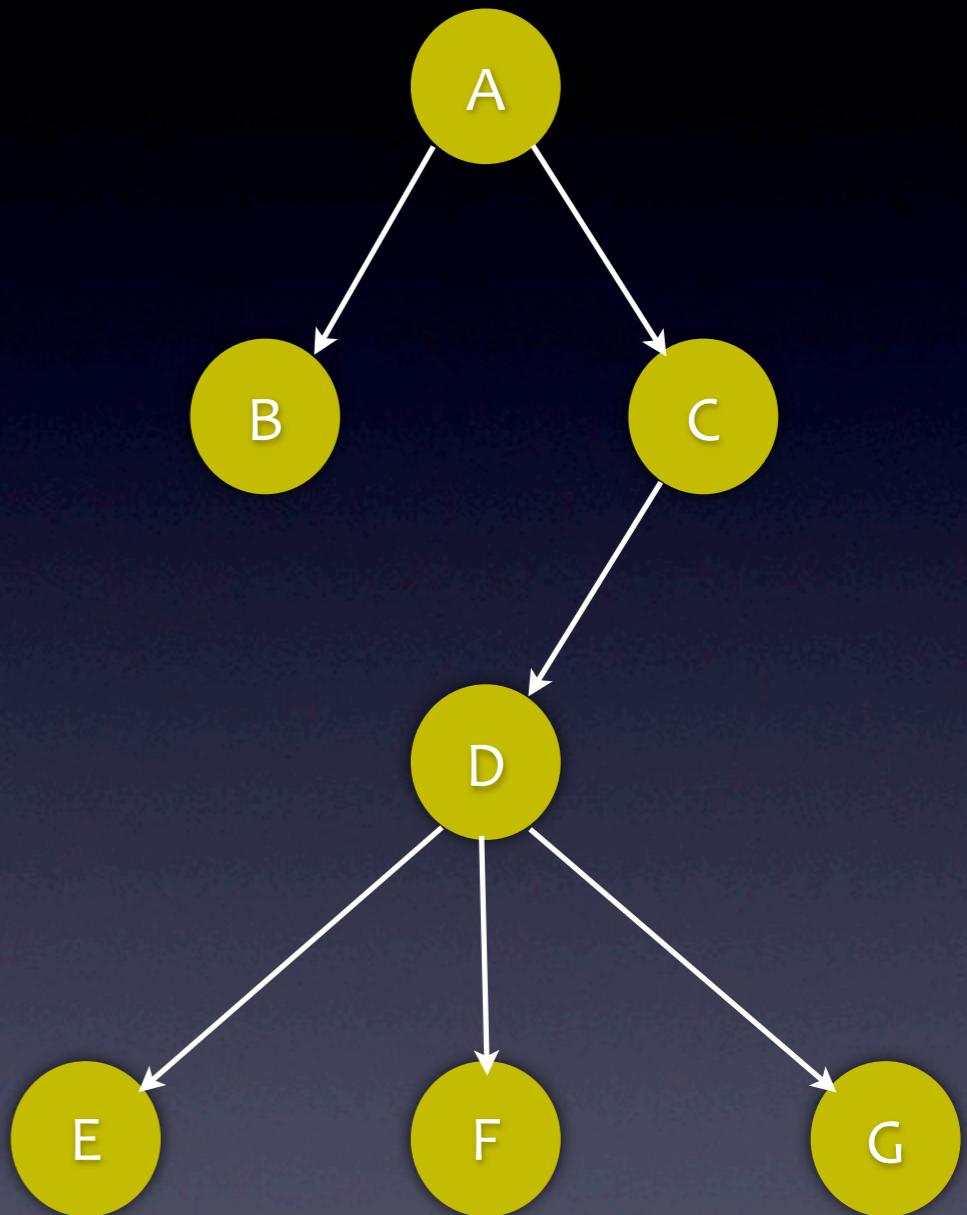
- Best practices based on several years of actor development
- Helpful hints for reasoning about actors at runtime

@jamie\_allen

何年かの開発に基づいたベスト・プラクティスとヒント

# Actors

- Concurrent, lightweight processes that communicate through asynchronous message passing
- Isolation of state, no internal concurrency



非同期なメッセージを用いて対話する並行プロセス  
状態の隔離

# Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case "Pong" => println("Pinging!"); sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case "Ping" => println("Ponging!"); sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Pong!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```

@jamie\_allen



# Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case "Pong" => println("Pinging!"); sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case "Ping" => println("Ponging!"); sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Pong!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```

@jamie\_allen



# Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case "Pong" => println("Ping!") ; sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case "Ping" => println("Pong!") ; sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Pong!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```

@jamie\_allen



# Akka Actors

```
class Pinger extends Actor {  
    def receive = {  
        case "Pong" => println("Ping!") ; sender ! "Ping!"  
    }  
}  
  
class Ponger extends Actor {  
    def receive = {  
        case "Ping" => println("Pong!") ; sender ! "Pong!"  
    }  
}  
  
object PingPong extends App {  
    val system = ActorSystem()  
    val pinger = system.actorOf(Props[Pinger])  
    val ponger = system.actorOf(Props[Ponger])  
    pinger.tell("Pong!", ponger)  
    Thread.sleep(1000)  
    system.shutdown  
}
```

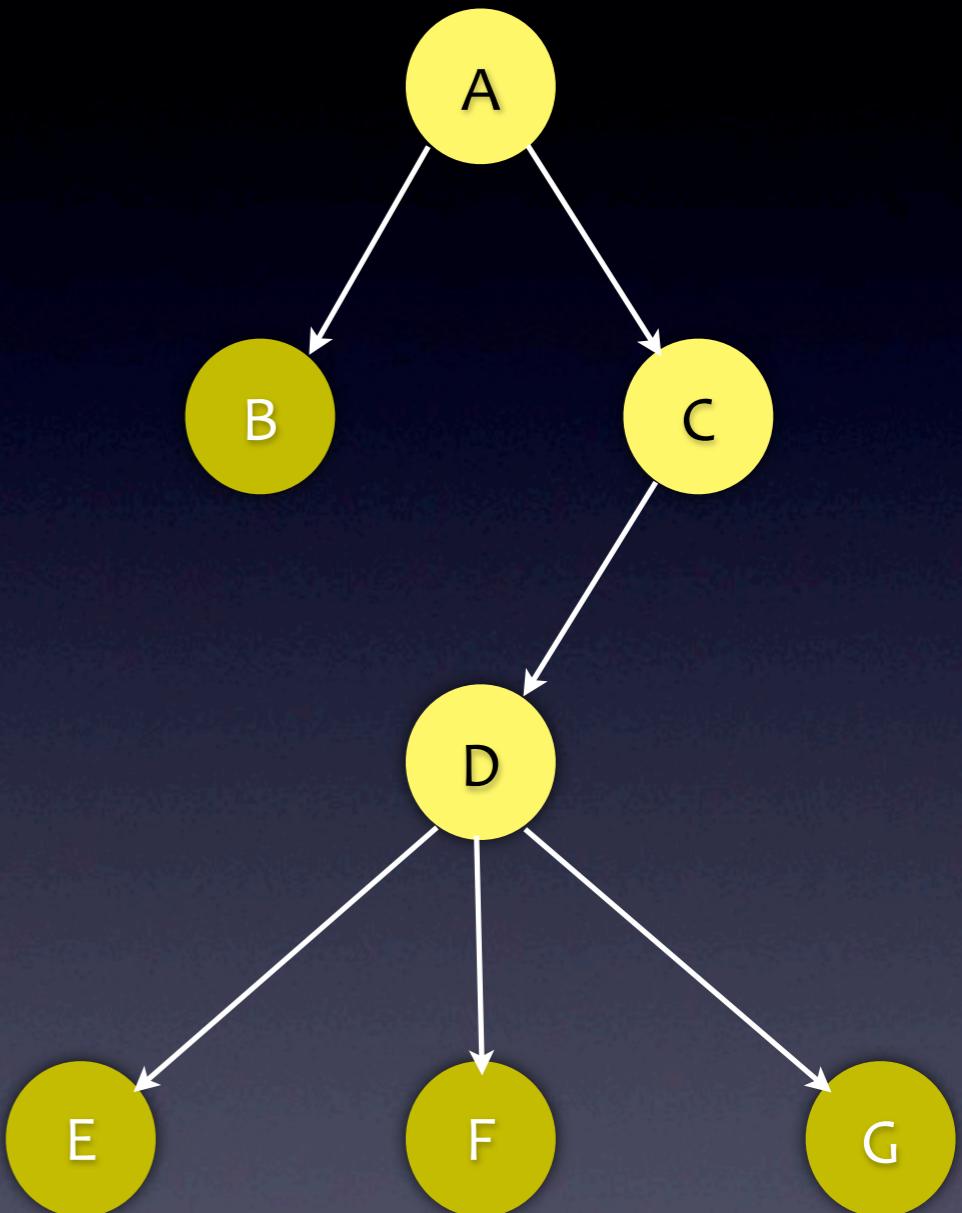
Overriding the “sender”

@jamie\_allen



# Supervisor Hierarchies

- Specifies handling mechanisms for groupings of actors in parent/child relationship



@jamie\_allen

Supervisor はアクター階層の処理を定義する

# Akka Supervisors

```
class MySupervisor extends Actor {  
    override val supervisorStrategy =  
        OneForOneStrategy() {  
            case ae: ArithmeticException => Resume  
            case np: NullPointerException => Restart  
        }  
  
    context.actorOf(Props[MyActor])  
}
```

@jamie\_allen



# Akka Supervisors

```
class MySupervisor extends Actor {  
    override val supervisorStrategy =  
        OneForOneStrategy() {  
            case ae: ArithmeticException => Resume  
            case np: NullPointerException => Restart  
        }  
  
    context.actorOf(Props[MyActor])  
}
```

@jamie\_allen



# Akka Supervisors

```
class MySupervisor extends Actor {  
    override val supervisorStrategy =  
        OneForOneStrategy() {  
            case ae: ArithmeticException => Resume  
            case np: NullPointerException => Restart  
        }  
  
    context.actorOf(Props[MyActor])  
}
```

Note the “context”



@jamie\_allen

# Domain Supervision

- Each supervisor manages a grouping of types in a domain
- Actors persist to represent existence of instances and contain their own state
- Actors constantly resolve the world as it should be against the world as it is

@jamie\_allen

Supervisor 階層を利用しドメインをモデルする

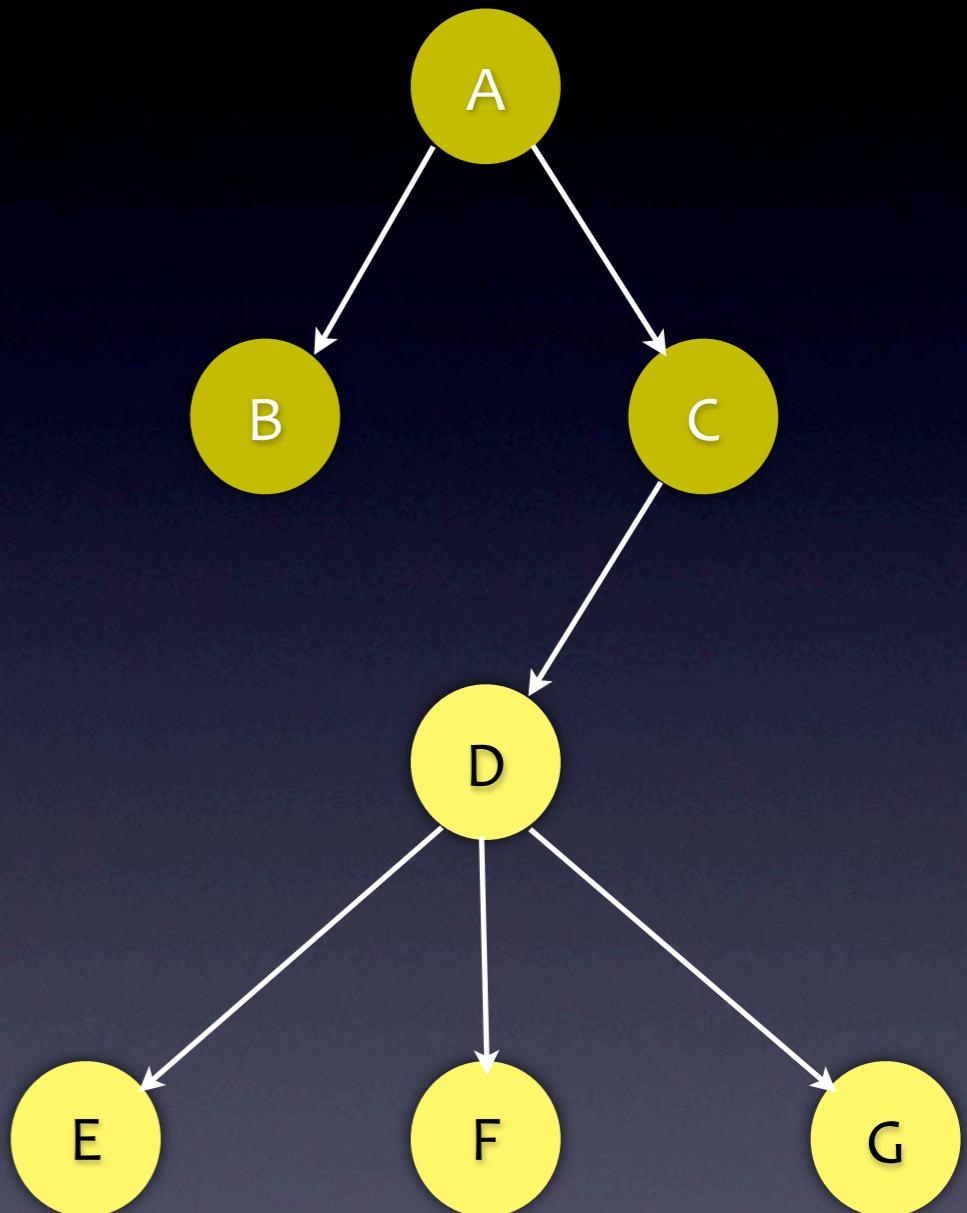
# Worker Supervision

- Supervisors should hold all critical data
- Workers should receive data for tasks in messages
- Workers being supervised should perform dangerous tasks
- Supervisor should know how to handle failures in workers in order to retry appropriately

Supervisor がデータを保持し、メッセージでワーカーに渡す  
ワーカーは死んでも構わない

# Parallelism

- Easily scale a task by creating multiple instances of an actor and applying work using various strategies
- Order is not guaranteed, nor should it be



@jamie\_allen

複数のアクターを作成することでタスクをスケール

# Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```

@jamie\_allen



# Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```

@jamie\_allen



# Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```

Should be configured externally



@jamie\_allen

# Akka Routing

```
class MyActor extends Actor {  
    def receive = { case x => println(x) }  
}  
  
object Parallelizer extends App {  
    val system = ActorSystem()  
    val router: ActorRef = system.actorOf(Props[MyActor] .  
        withRouter(RoundRobinRouter(nrOfInstances = 5)))  
  
    for (i <- 1 to 10) router ! i  
}
```

@jamie\_allen



# RULE: Actors Should Only Do One Thing

@jamie\_allen

アクターは1つの事に特化すること

# Single Responsibility Principle

- Do not conflate responsibilities in actors
- Becomes hard to define the boundaries of responsibility
- Supervision becomes more difficult as you handle more possibilities
- Debugging becomes very difficult

@jamie\_allen

单一責任の原則: なんでもかんでも手を出さない

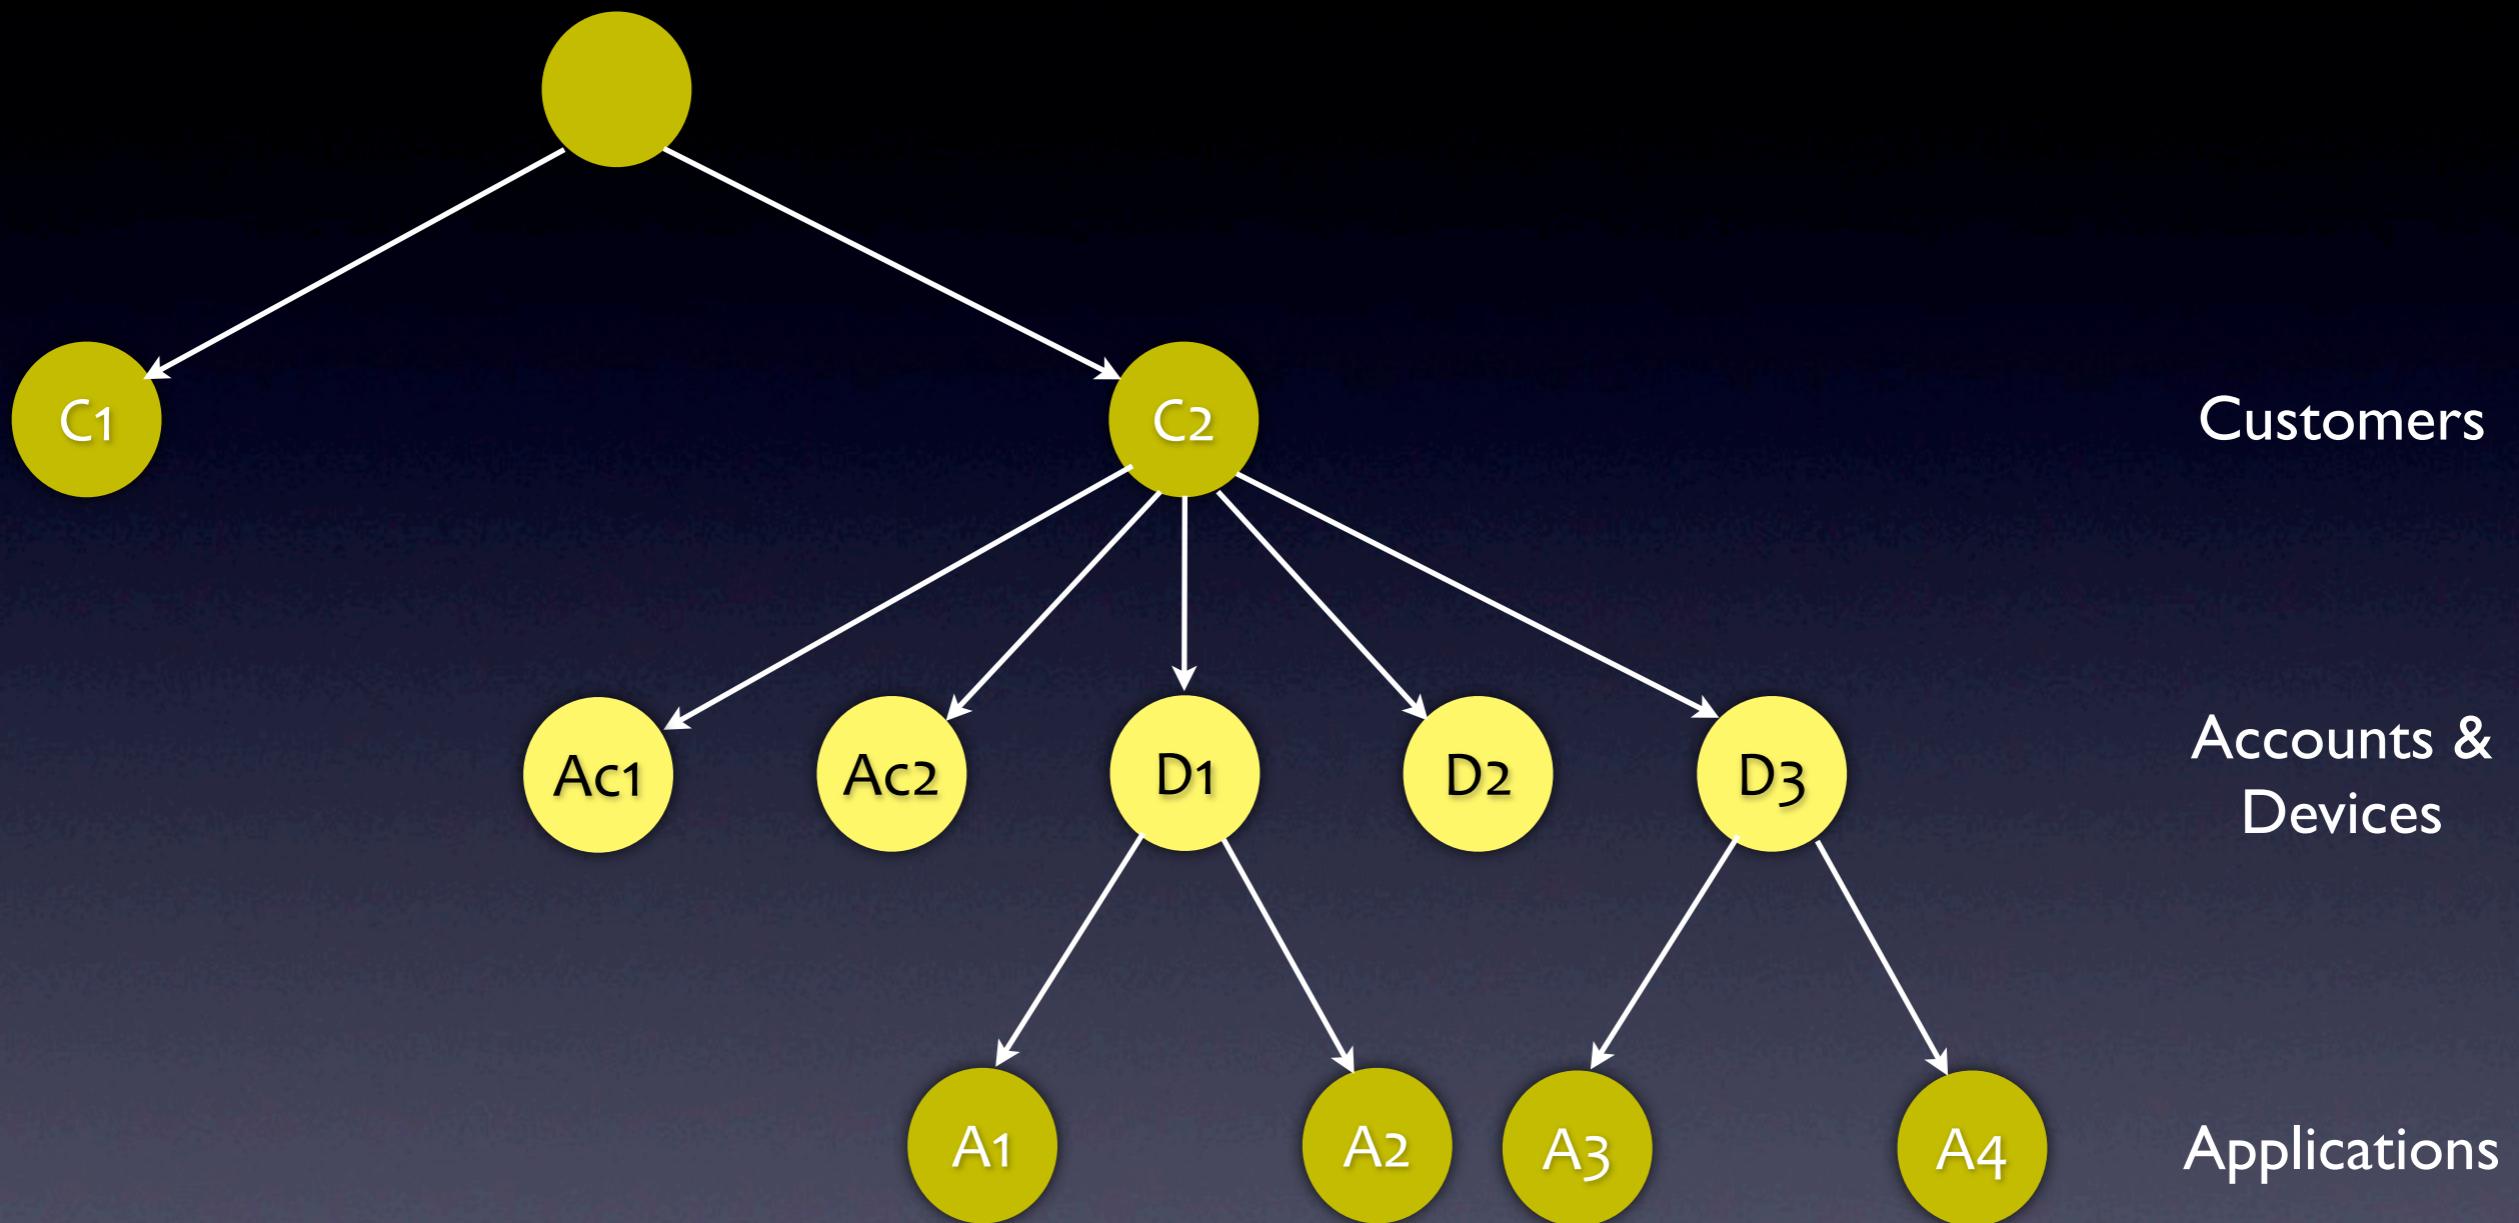
# Supervision

- Every non-leaf node is technically a supervisor
- Create explicit supervisors under each node for each type of child to be managed

@jamie\_allen

アクターの種類ごとに Supervisor を分けること

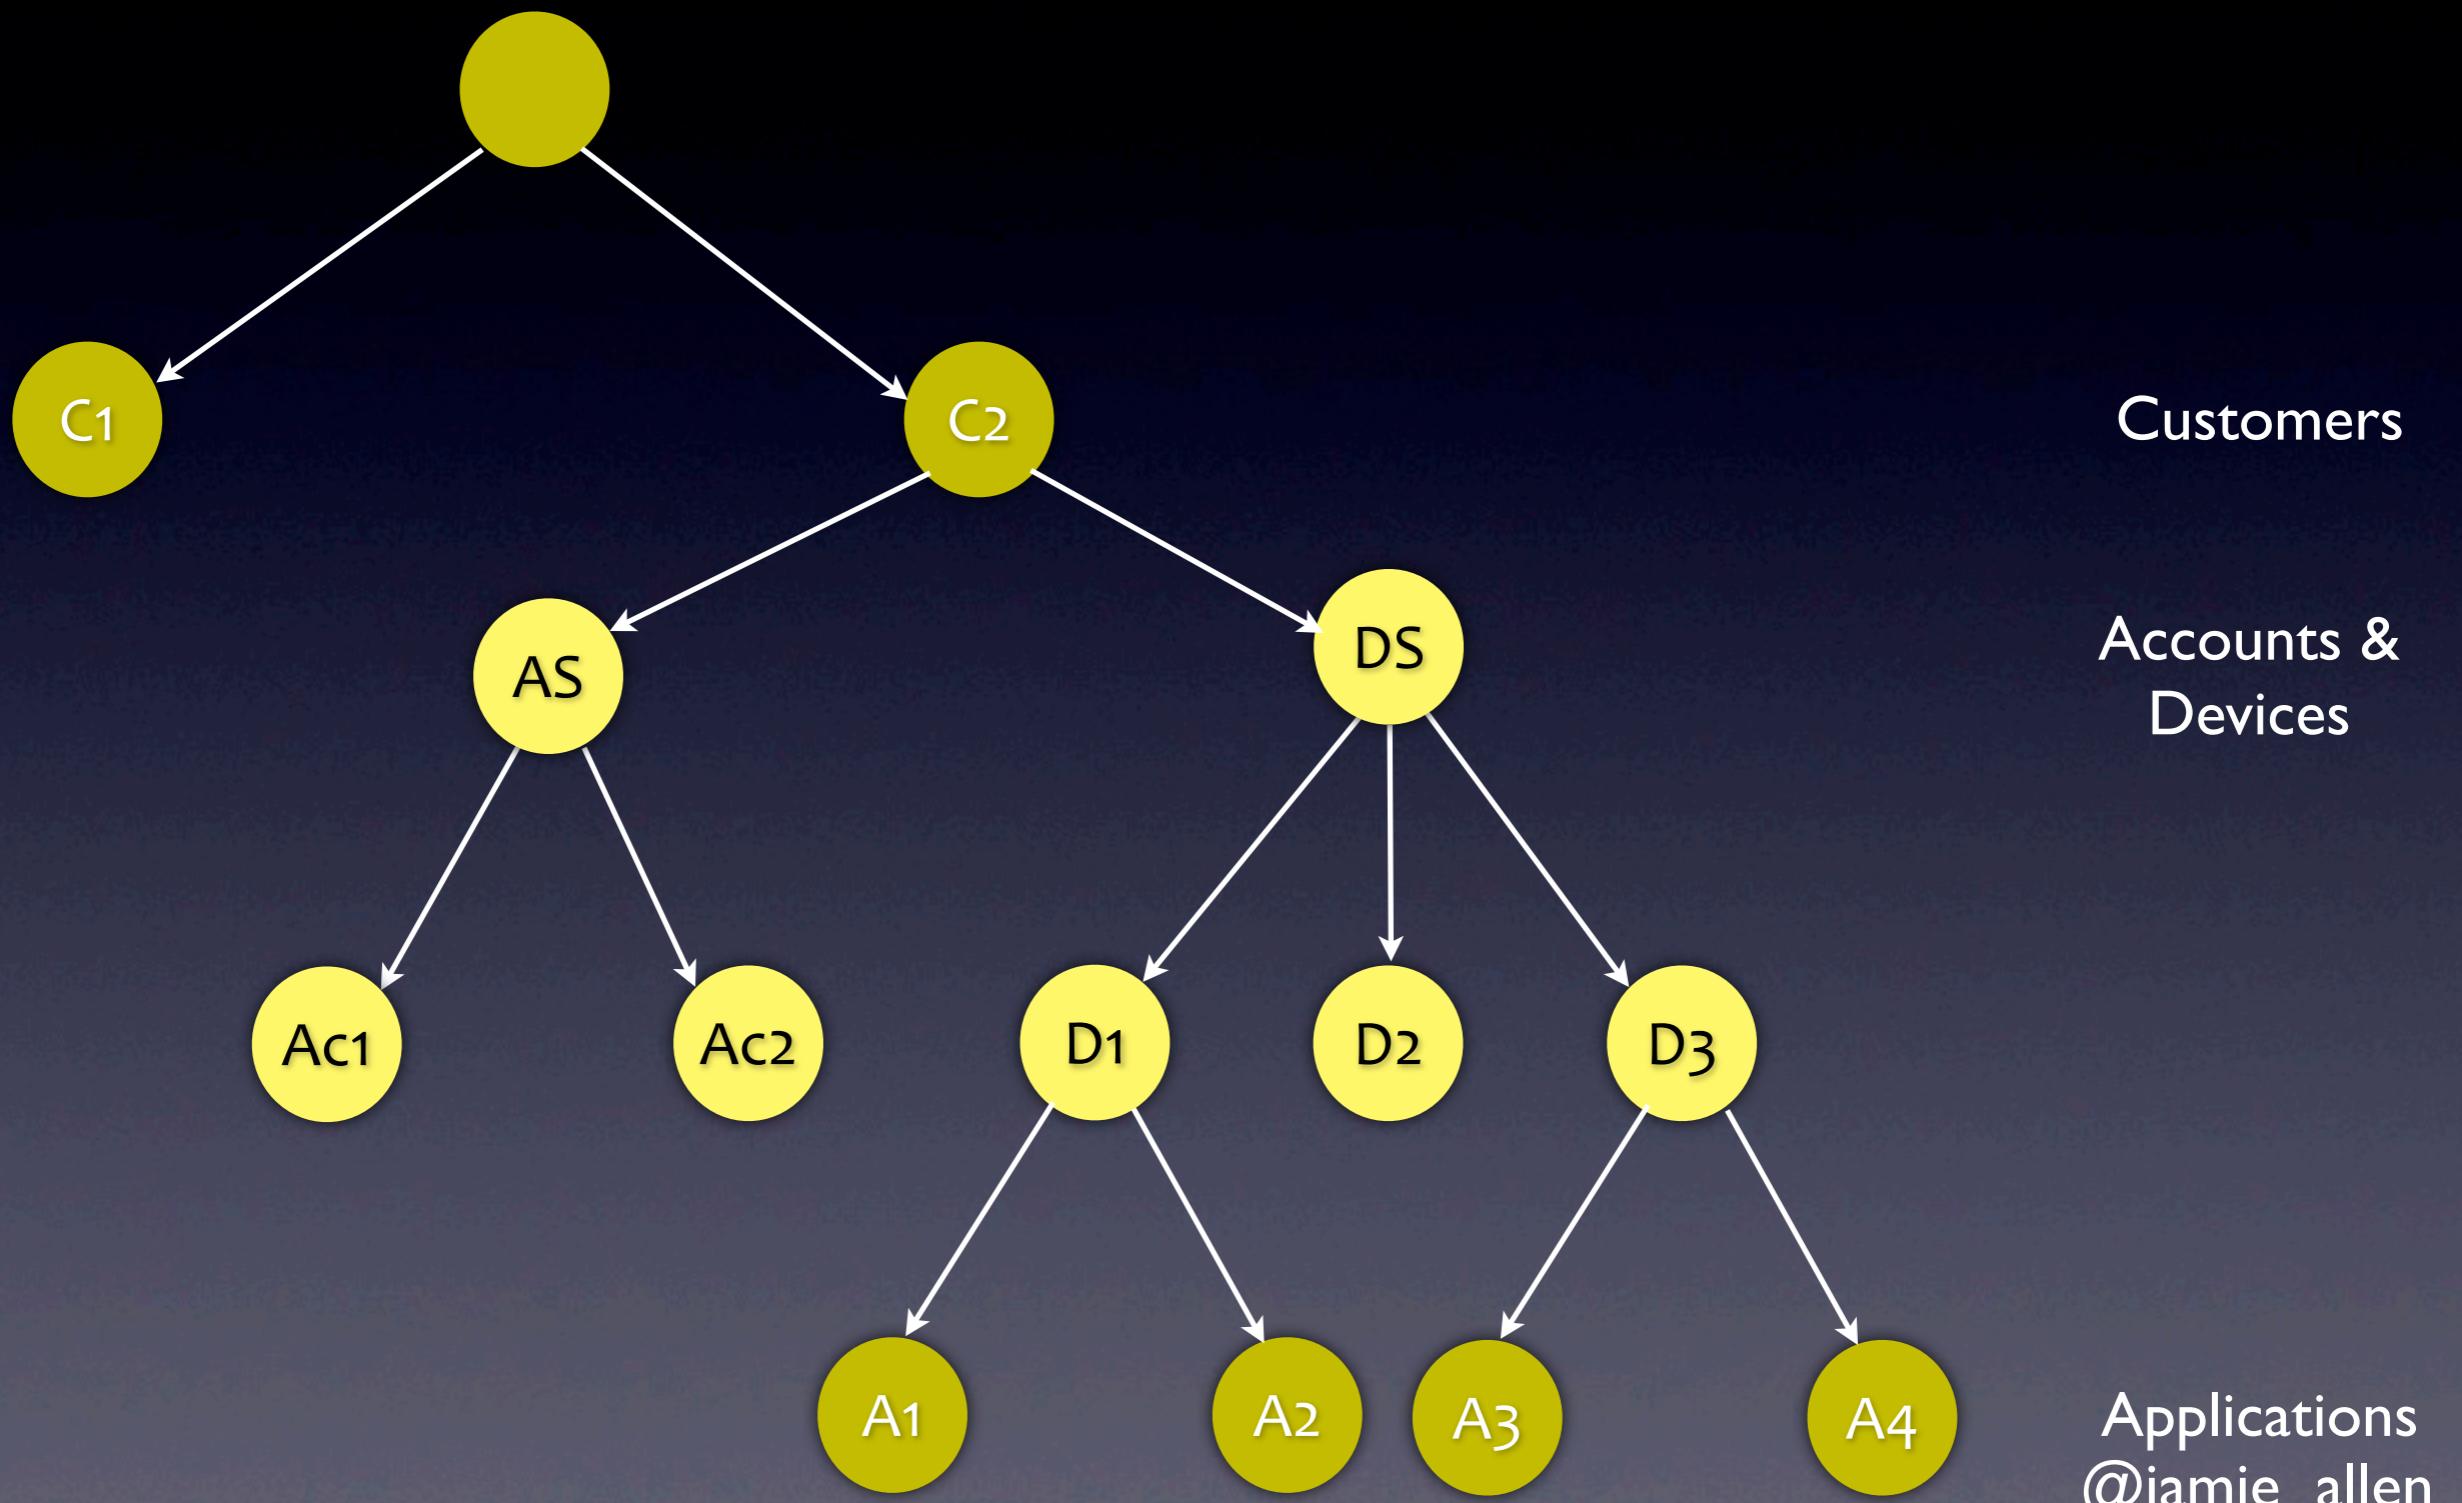
# Conflated Supervision



@jamie\_allen

混ざってる Supervisor

# Explicit Supervision



@jamie\_allen

明示的な Supervisor

# Keep the Error Kernel Simple

- Limit the number of supervisors you create at this level
- Helps with fault tolerance and explicit handling of errors through the hierarchy
- Akka uses synchronous messaging to create top-level actors

@jamie\_allen

Error Kernel はシンプルに

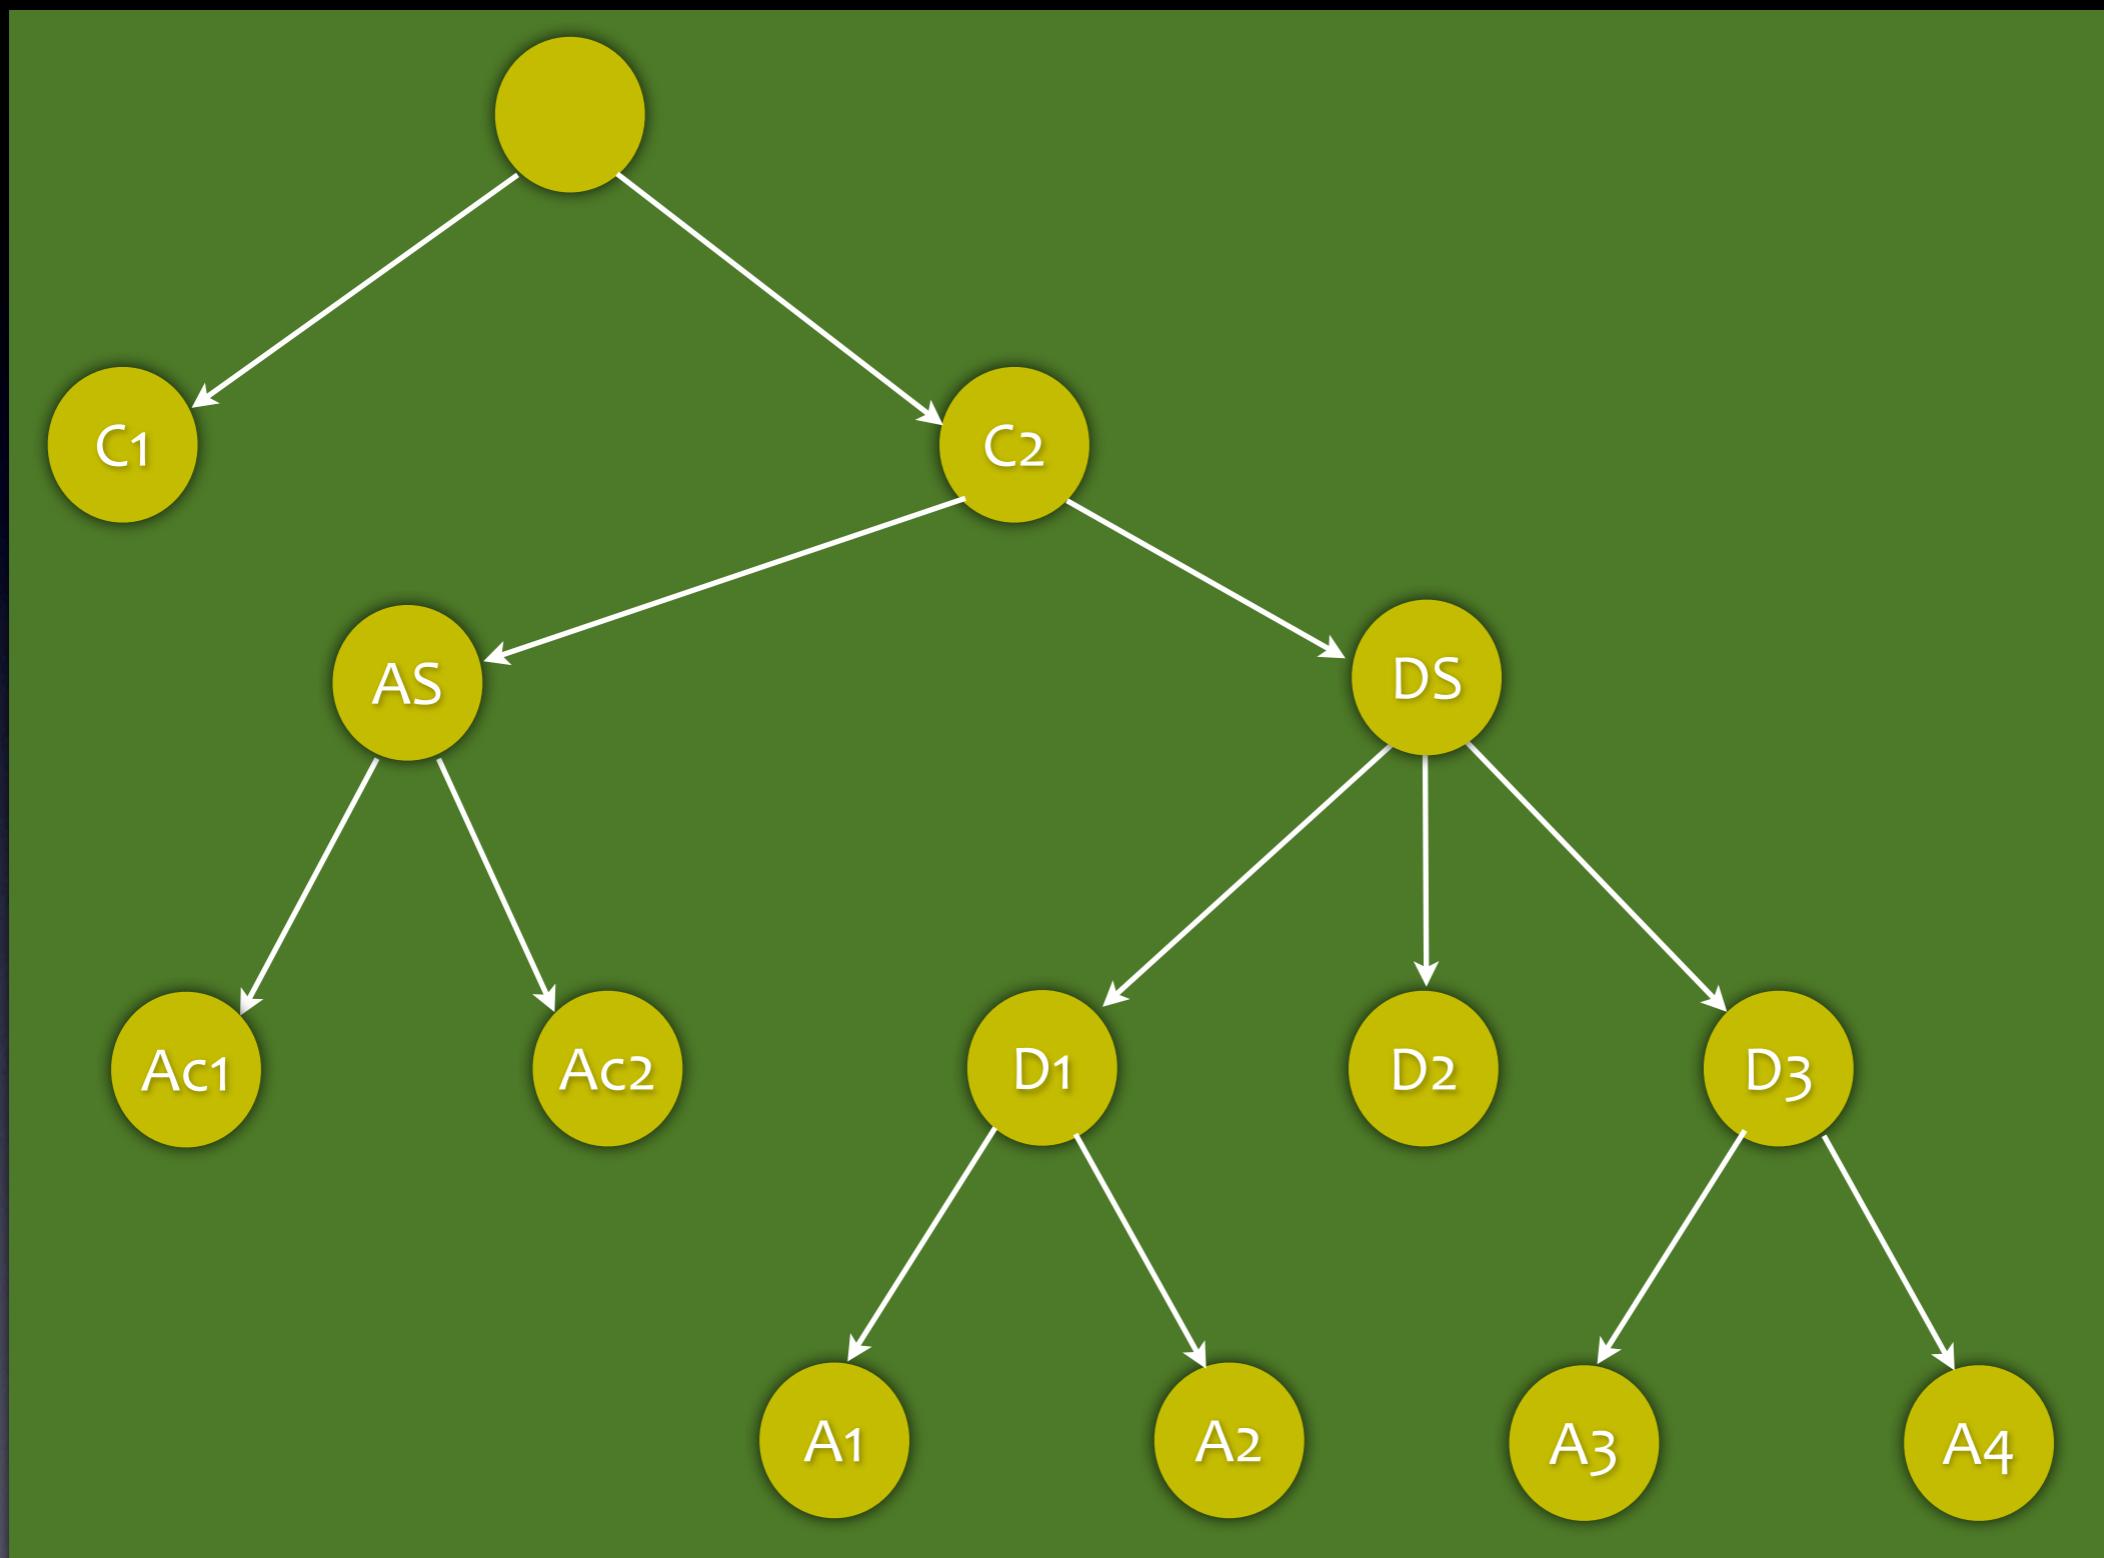
# Use Failure Zones

- Multiple isolated zones with their own resources (thread pools, etc)
- Prevents starvation of actors
- Prevents issues in one branch from affecting another

@jamie\_allen

Failure Zone を使ってリソースを隔離する

# Failure Zones



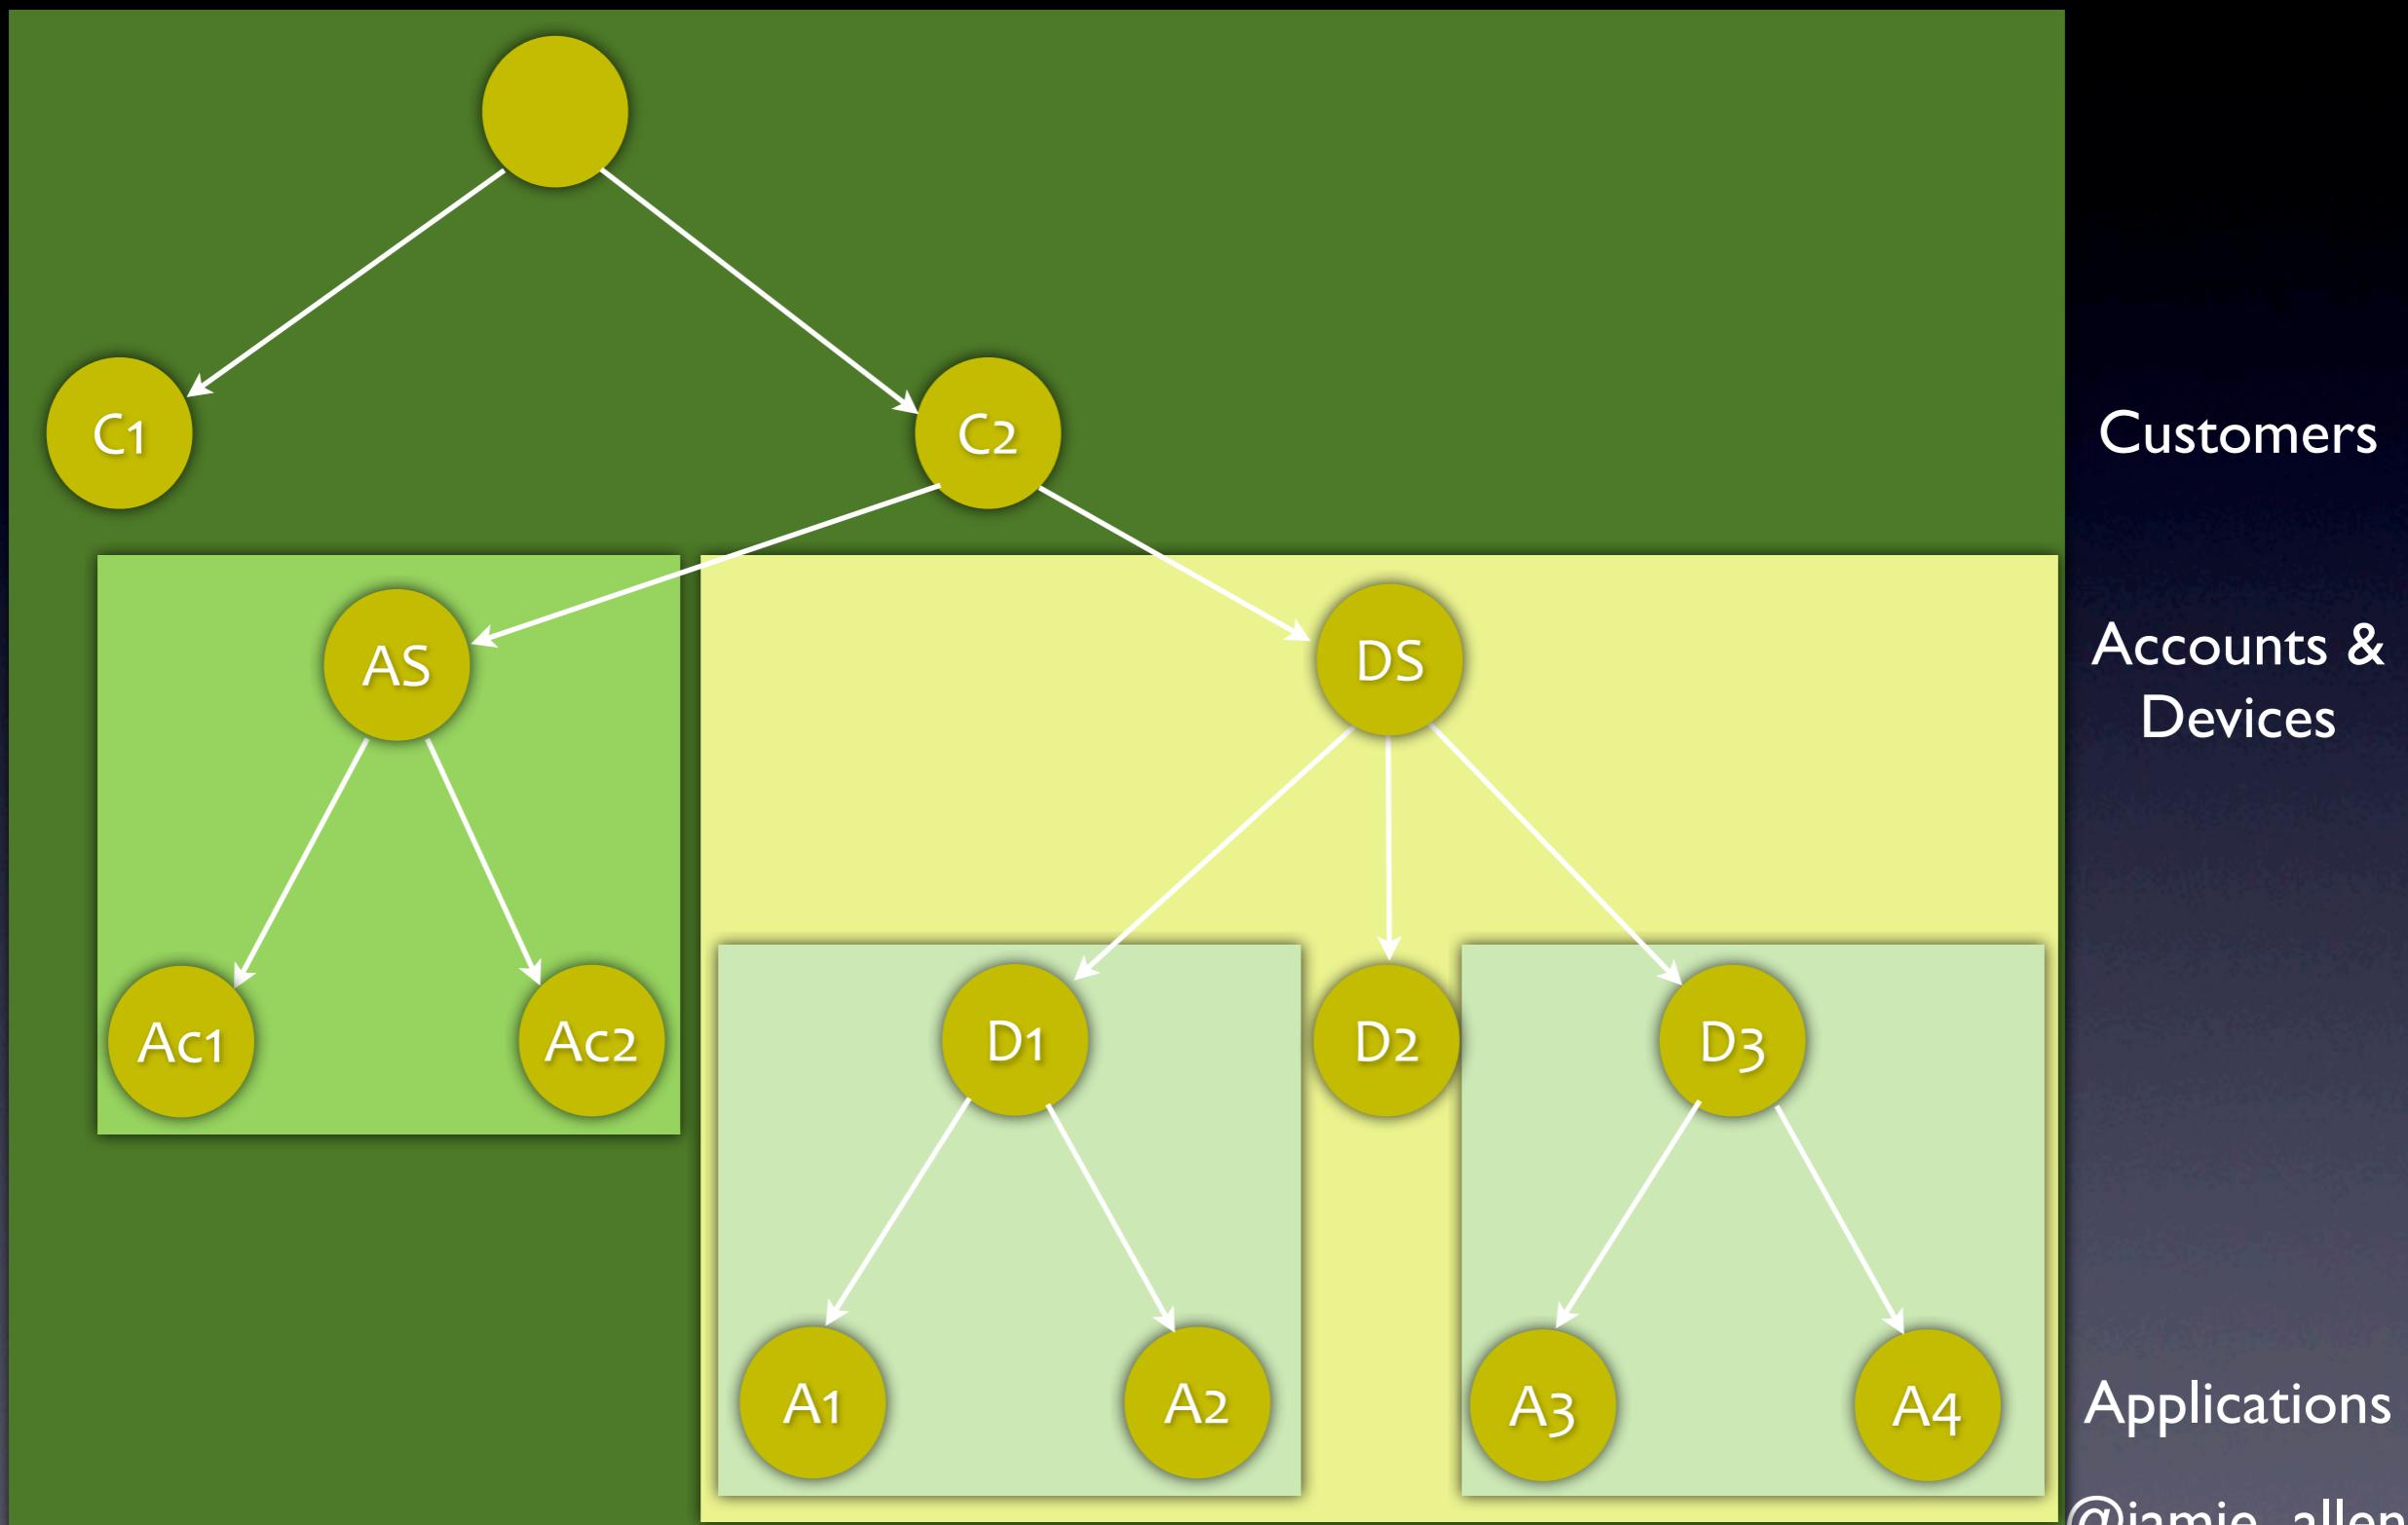
Customers

Accounts &  
Devices

Applications

@jamie\_allen

# Failure Zones



@jamie\_allen

# Takeaway

- Isolation of groups of actors limits the effects of failure
- For reasonably complex actor systems, shallow trees are a smell test
- Actors are cheap - use them

@jamie\_allen

アクターのグループを隔離することで障害の影響を抑える

# RULE: Block Only When and Where You Must

@jamie\_allen

必要な時に、必要な場所のみでブロックすること

# Consequences of Blocking

- Eventually results in actor starvation as thread pool dries up
- Horrible performance
- Massive waste of system resources

ブロッキングの行く末: 無くなるスレッド、飢えるアクター、  
落ちる性能

# Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case "Start" =>  
            worker ? (1 to 100) map  
                (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```

@jamie\_allen



# Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case "Start" =>  
            worker ? (1 to 100) map  
                (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```

@jamie\_allen



# Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case "Start" =>  
            worker ? (1 to 100) map  
                (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```

@jamie\_allen



# Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case "Start" =>  
            worker ? (1 to 100) map  
                (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```

@jamie\_allen



# Futures

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker])  
    def receive = {  
        case "Start" =>  
            worker ? (1 to 100) map  
                (x => println("Got value: " + x))  
    }  
}  
  
object Bootstrapper extends App {  
    val system = ActorSystem()  
    val delegator = system.actorOf(Props[Delegator])  
    delegator ! "Start"  
    Thread.sleep(1000)  
    system.shutdown  
}
```

@jamie\_allen



# Sequential Futures

```
val r: Future[Int] = for {
  a <- (service1 ? GetResult).mapTo[Int]
  b <- (service2 ? GetResult(a)).mapTo[Int]
} yield a * b
```

@jamie\_allen

順次 Future

# Sequential Futures

Type Safety

```
val r: Future[Int] = for {  
    a <- (service1 ? GetResult).mapTo[Int]  
    b <- (service2 ? GetResult(a)).mapTo[Int]  
} yield a * b
```

@jamie\_allen

順次 Future

# Parallel Futures

## First way: Define futures first

```
val f1: Future[Int] = service1 ? GetResult  
val f2: Future[Int] = service2 ? GetResult  
  
val r: Future[Int] = for {  
    a <- f1  
    b <- f2  
} yield a * b
```

@jamie\_allen

並列 Future 方法1: 事前に Future を定義

# Parallel Futures

Second way: Define in one line of for comprehension

```
val r: Future[Int] = for {  
    (a: Int, b: Int) <- (service1 ? GetResult) zip (service2 ? GetResult)  
} yield a * b
```



Note the “zip”

@jamie\_allen

並列 Future 方法2: zip を使って 1行で定義

# Blocking

- An example is database access
- Use a specialized actor with its own resources
- Pass messages to other actors to handle the result
- Scala provides “Managed Blocking” to limit the number of blocking operations taking place in actors at one time

@jamie\_allen

ブロッキング専門のアクターを作る

# Akka Dispatcher

```
class Worker extends Actor {  
    def receive = {  
        case s: Seq[Int] => sender ! s.reduce(_ + _)  
    }  
}  
  
class Delegator extends Actor {  
    implicit val timeout: Timeout = 2 seconds  
    val worker = context.actorOf(Props[Worker]).withDispatcher("my-dispatcher") ←———— Failure Zone  
    def receive = {  
        case _ =>  
            blocking {  
                val futResult = worker ? (1 to 100)  
                val result = Await.result(futResult, 2 seconds)  
            }  
    }  
}
```

@jamie\_allen



# Push, Not Pull

- Start with no guarantees about delivery
- Add guarantees only where you need them
- Retry until you get the answer you expect
- Switch your actor to a "nominal" state at that point

@jamie\_allen

配達時間はまず保証無しから始める

# Takeaway

- Find ways to ensure that your actors remain asynchronous and non-blocking
- Avoid making your actors wait for anything while handling a message

@jamie\_allen

アクターが非同期で非ブロッキングであるように努めること

# RULE: Do Not Optimize Prematurely

@jamie\_allen

早すぎる最適化はダメ

# Start Simple

- Make Donald Knuth happy
- Start with a simple configuration and profile
- Do not parallelize until you know you need to and where

@jamie\_allen

プロファイルも取らずに並列化しないこと

# Initial Focus

- Deterministic
- Declarative
- Immutable
- Start with functional programming and go from there

@jamie\_allen

まずは決定的、宣言型、不变など関数型の基本から

# Advice From Jonas Bonér

- Layer in complexity
- Add indeterminism
- Add mutability in hot spots  
(CAS and STM)
- Add explicit locking and  
threads as a last resort



Photo courtesy of Brian Clapper, NE Scala 2011

段階的に複雑さを導入する: 1. 非決定性 2. 可変性  
3. スレッドと明示的なロックは最終手段

# Prepare for Race Conditions

- Write actor code to be agnostic of time and order
- Actors should only care about now, not that something happened before it
- Actors can "become" or represent state machines to represent transitions

@jamie\_allen

アクターは時間と順序に無関心であるべき

# Beware the Thundering Herd

- Actor systems can be overwhelmed by "storms" of messages flying about
- Do not pass generic messages that apply to many actors
- Dampen actor messages if the exact same message is being handled repeatedly within a certain timeframe
- Tune your dispatchers and mailboxes via back-off policies and queue sizes
- Akka now has Circuit Breakers

メッセージの「嵐」に注意

複数のアクターに作用する汎用メッセージは避ける

# Takeaway

- Start by thinking in terms of an implementation that is deterministic and not actor based
- Layer in complexity as you go

@jamie\_allen

まずはアクターを使わない決定的な実装から

# RULE: Be Explicit In Your Intent

@jamie\_allen

意図をハッキリさせること

# Name Your Actors

- Allows for external configuration
- Allows for lookup
- Better semantic logging

```
val system = ActorSystem("pingpong")
val pinger = system.actorOf(Props[Pinger], "pinger")
val ponger = system.actorOf(Props[Ponger], "ponger")
```

@jamie\_allen

アクターには名前を付けること

# Create Specialized Messages

- Non-specific messages about general events are dangerous

`AccountsUpdated`

- Can result in "event storms" as all actors react to them
- Use specific messages forwarded to actors for handling

`AccountDeviceAdded (acctNum, deviceNum)`

@jamie\_allen

専用メッセージを作る。汎用イベントは危険。

# Create Specialized Exceptions

- Don't use `java.lang.Exception` to represent failure in an actor
- Specific exceptions can be handled explicitly
- State can be transferred between actor incarnations in Akka (if need be)

@jamie\_allen

専用例外を作る

# Takeaway

- Be specific in everything you do
- Makes everything that occurs in your actor system more clear to other developers maintaining the code
- Makes everything more clear in production

@jamie\_allen

全てにおいて特定であること。本番環境でも役立つ。

# RULE: Do Not Expose Your Actors

@jamie\_allen

アクターを公開しないこと

# No Direct References

- Actors die
- Doesn't prevent someone from calling into an actor with another thread
- Akka solves this with the ActorRef abstraction
- Erlang solves this with PIDs

@jamie\_allen

絶対にアクターを直接参照してはいけない

# Never Publish “this”

- Don't send it anywhere
- Don't register it anywhere
- Particularly with future callbacks
- Publish “self” instead, which is an ActorRef
- Avoid closing over "sender" in Akka, it will change with the next message

@jamie\_allen

絶対に this を返してはいけない。代わりに self を使う。

# Use Immutable Messages

- Enforces which actor owns the data
- If mutable state can escape, what is the point of using an actor?

@jamie\_allen

不变メッセージを使うこと

# Pass Copies of Mutable Data

- Mutable data in actors is fine
- But data can escape your scope
- Copy the data and pass that, as Erlang does (COW)
- Akka has STM references

@jamie\_allen

可変データはコピーを送る

# Avoid Sending Behavior

- Unless using Agents, of course
- Closures make this possible (and easy)
- Also makes it easy for state to escape

@jamie\_allen

振る舞いを送ることは避ける

# Takeaway

- Keep everything about an actor internal to that actor
- Be vary wary of data passed in closures to anyone else

@jamie\_allen

アクターの内部構造を外に見せないこと

# RULE: Make Debugging Easy On Yourself

@jamie\_allen

デバッグが簡単にできるようにする

# Externalize Business Logic

- Consider using external functions to encapsulate complex business logic
- Easier to unit test outside of actor context
- Not a rule of thumb, but something to consider as complexity increases
- Not as big of an issue with Akka's TestKit

@jamie\_allen

ビジネス・ロジックは外に書く

# Use Semantically Useful Logging

- Trace-level logs should have output that you can read easily
- Use line-breaks and indentation
- Both Akka and Erlang support hooking in multiple listeners to the event log stream

@jamie\_allen

読んで意味のあるログを書く

# Unique IDs for Messages

- Allows you to track message flow
- When you find a problem, get the ID of the message that led to it
- Use the ID to grep your logs and display output just for that message flow
- Akka ensures ordering on a per actor basis, also in logging

@jamie\_allen

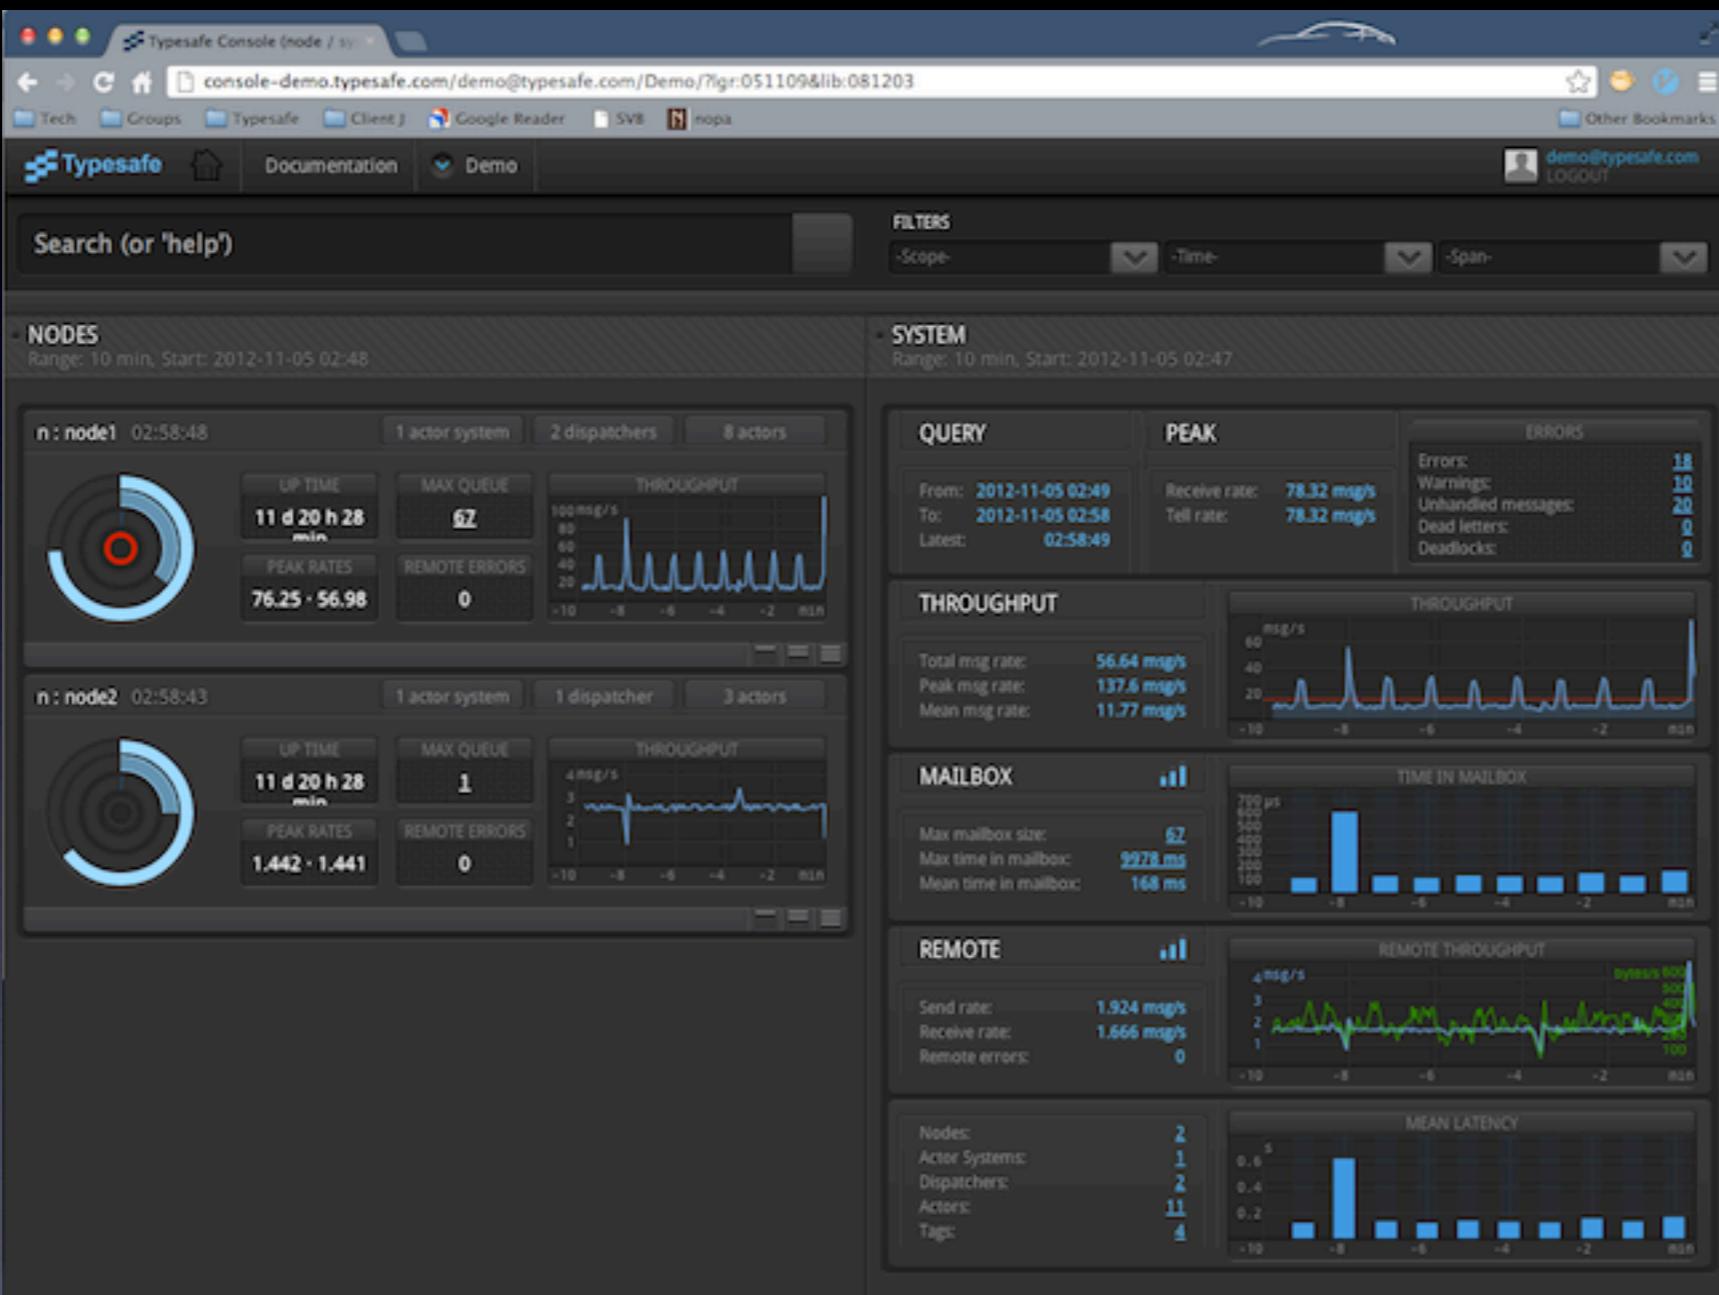
メッセージにユニークIDをふる

# Monitor Everything

- Do it from the start
- Use tools like JMX MBeans to visualize actor realization
- The Typesafe Console is a great tool to visualize actor systems, doesn't require you to do anything up front
- Visual representations of actor systems at runtime are invaluable

@jamie\_allen

# Typesafe Console



@jamie\_allen

# Takeaway

- Build your actor system to be maintainable from the outset
- Utilize all of the tools at your disposal

@jamie\_allen

メンテが容易なアクターシステムを構築する

# Thank You!

- Some content provided by members of the Typesafe team, including:
  - Jonas Bonér
  - Viktor Klang
  - Roland Kuhn
  - Havoc Pennington

@jamie\_allen

ご清聴ありがとうございました